Last Time

α vectors

$\alpha_1[s] \cdot [0.4, 0.3, 0.2]$

set Γ

$$V(b) = \max_{\alpha \in \Gamma} \alpha^T b$$

$S = \{1, 2\}$

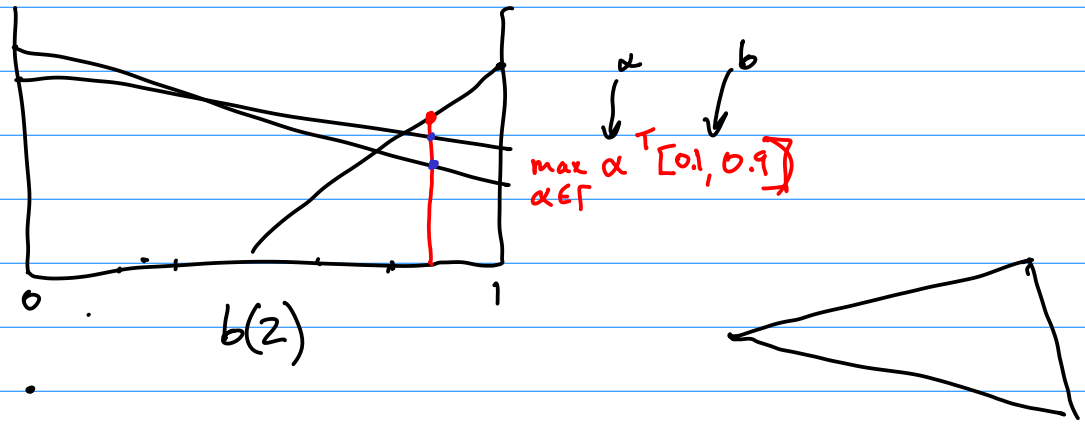

Today

Offline POMDP

Name	S	A	O	
Exact VI	D (10)	D	D	
PBVI	D	D	D	bigger
Perseus	D	D	D	
HSVI	D	D	D	
SARSOP	D (1000)	D	D	
MCVI	C	D	D/C	
LQG	C	C	C	Linear Gaussian Quadratic

"Survey of Point-Based POMDP Solvers"

Exact VI

loop

$$\Gamma' \doteq \bigcup_{a \in A} \Gamma^a$$

$$\Gamma^a = \bigoplus_{o \in O} \Gamma^{ao}$$

$$\Gamma^{ao} = \left\{ \frac{1}{|O|} r_a + \alpha^{a,o} : \alpha \in \Gamma \right\}$$

$$\alpha^{a,o}[s] = \sum_{s'} Z(o|a,s') T(s'|s,a) \alpha[s']$$

for 1 iteration

$$O(|\Gamma||A||O||S|^2 + |A||S||\Gamma|^{|O|})$$

PBVI

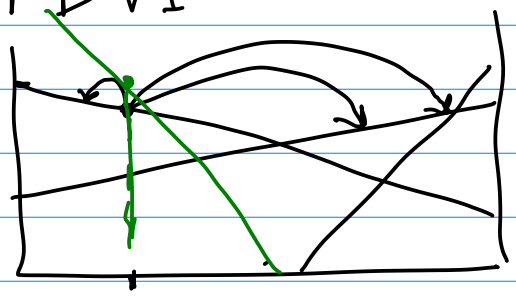

→ Backup Belief(Γ, b)

for $a \in A$

for $o \in O$

$b' = \tau(b, a, o)$

$\alpha_{a,o} = \underset{\alpha \in \Gamma}{\text{argmax}}\ \alpha^T b'$

for $s \in S$

$\alpha_a[s] = R(s,a) + \gamma \sum_{s',o} T(s'|s,a) Z(o'|a,s') \alpha_{a,o}[s']$

return $\underset{\alpha_a}{\text{argmax}}\ \alpha_a^T b$

$O(|\Gamma||A||O||S|^2 + |A||S||\Gamma||B|)$ ← set of Beliefs

How we choose B

Original PBVI

$B = \{b_0\}$

loop $B' = \phi$

for $b \in B$

$\Gamma = \Gamma \cup$ Backup Belief (Γ, b)

for $b \in B$

$\tilde{B} = \{\tau(b,a,o) : a \in A, o \in O\}$

$B' = B' \cup \underset{b \in \tilde{B}}{\text{argmax}} \|B, b\|$ ← add furthest b away from B

$B = B \cup B'$

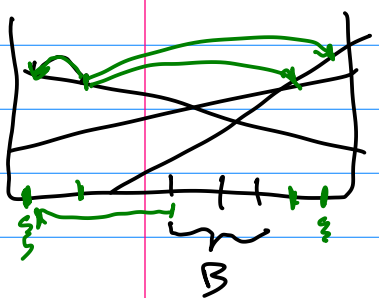

$B = \{b^1, b^2, b^3\}$

$B' = \{b^4, b^5\}$

$B \cup B' = \{b^1 \dots b^5\}$

Perseus - randomly choose B

$B = \emptyset$

$b = b_0$

loop until $|B| = n$

$a = rand(A)$

$o = rand(P(o|b,a))$

$B = B \cup \{\tau(b,a,o)\}$

HSVI Heuristic Search Value Iteration

$\bar{V}(b)$ upper bound

$\underline{V}(b)$ lower bound

while $\bar{V}(b_0) - \underline{V}(b_0) > \varepsilon$

explore $(b_0, 0)$

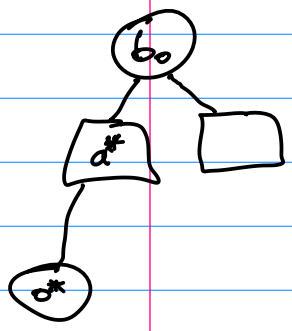

explore (b,t)

if $\bar{V}(b) - \underline{V}(b) > \varepsilon \gamma^t$

$a^* = \underset{a}{\text{argmax}}\ \bar{Q}(b,a)$ ← upper bound

$o^* = \underset{o}{\text{argmax}} \left(P(o|b,a) \left(\bar{V}(\tau(b,a,o)) - \underline{V}(\tau(b,a,o)) - \varepsilon \gamma^{(t+1)} \right) \right)$ ← excess uncertainty

explore $(\tau(b,a^*,o^*), t+1)$

$\underline{\Gamma} = \underline{\Gamma} \cup$ Backup Belief $(\underline{\Gamma}, b)$

$\bar{V}(b) = B_b[\bar{V}(b)]$

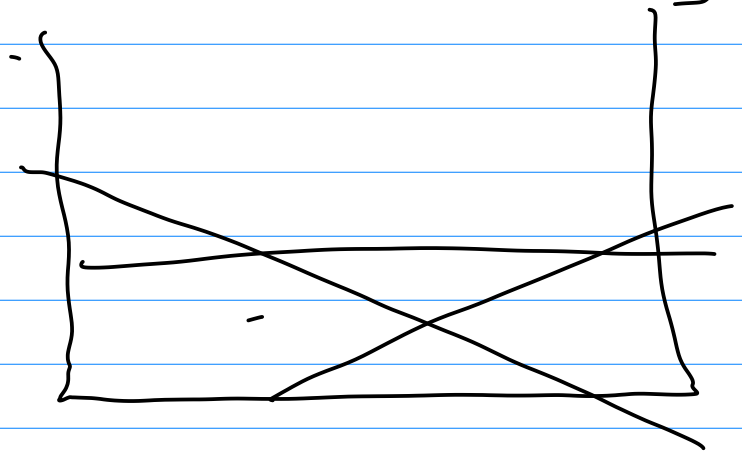

Baby POMDP

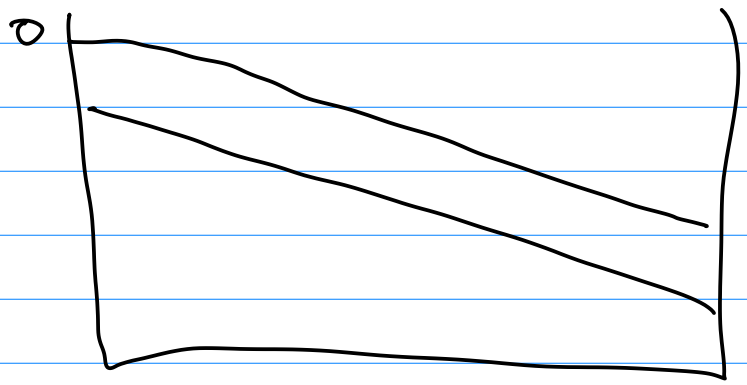

$\frac{R}{1-\gamma}$ Open-Loop Policy

Never Feed

$\frac{-5}{1-0.9} = -50$

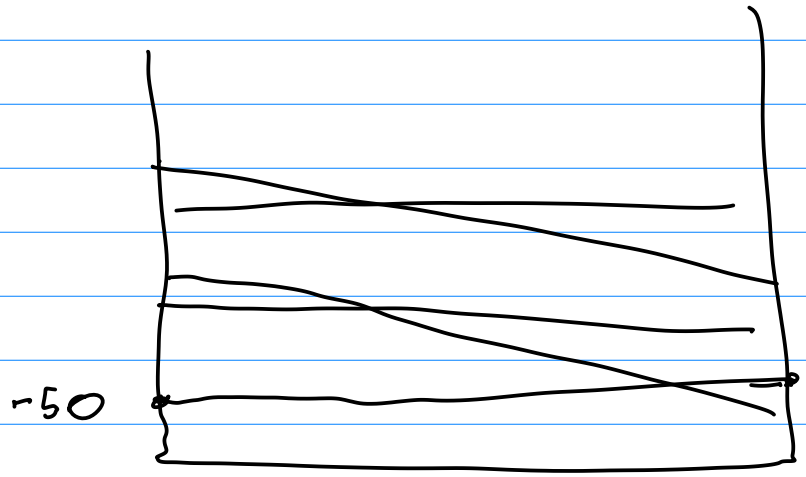

Upper Bounds

Sawtooth Upper bound

Finding points to interpolate between is an LP

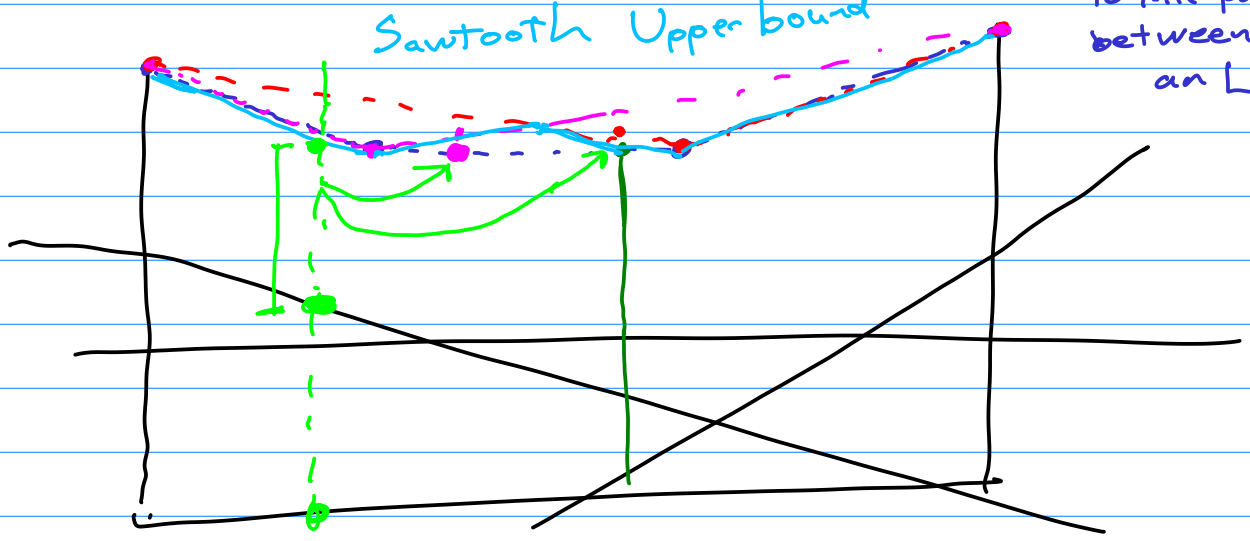

$$B_b[\bar{V}(b)] = \max R(b,a) + \gamma \sum_o P(o|b,a)\bar{V}(\tau(b,a,o))$$

SARSOP — Successive Approximation of Reachable Space under Optimal Policies

Similar to HSVI

BCR BCR^*

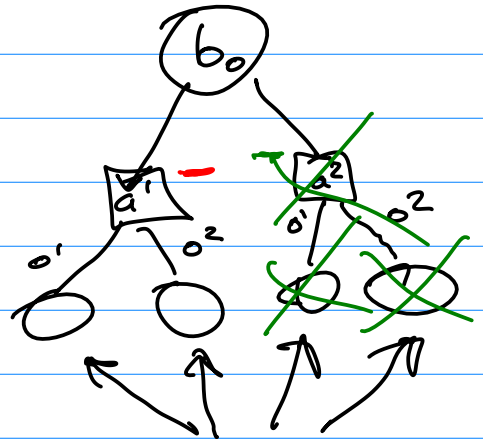

Belief points that could be in B

if $\bar{Q}(b,a^1) < \underline{Q}(b,a^2)$
then prune all b below (b,a^2)
from B

Instead of pruning α that are dominated over the whole belief space, prune α dominated over B

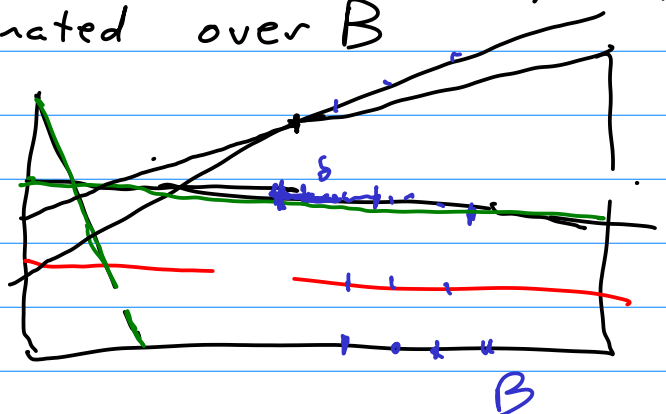

Pruned by Any

Pruned by SARSOP

Policy Graph

vertex labeled with a
edge labeled with o

Feed when Crying

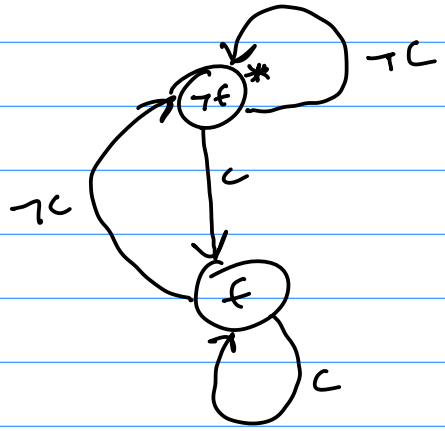

Policy Graph → α vector
evaluate plan
α vectors → policy graphs

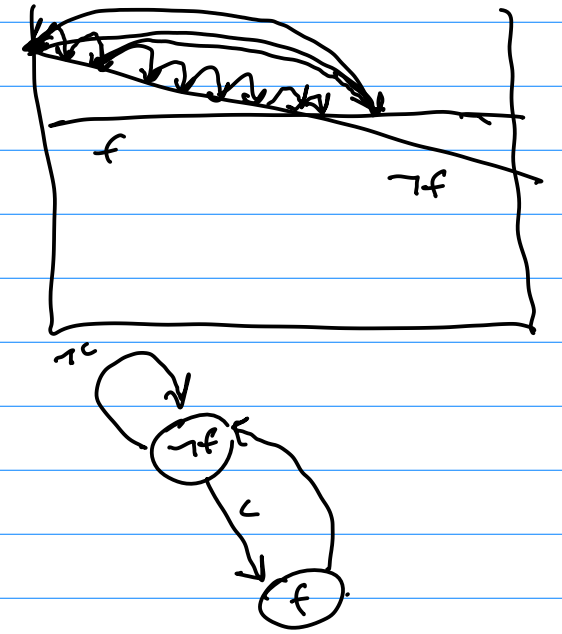

MCVI "Monte Carlo Value Iteration"

$$V_G(b) = \max \int_{s \in S} \alpha_v(s)\, b(s)\, ds$$ ← Approx w/MC

MC.

v

MC-Backup (G, b, N)

$R_a = 0 \quad V_{a,o,v} = 0$

for $a \in A$

for i in $1:N$

$s_i \Leftarrow$ sample (b)

$s_i', o_i, r_i \Leftarrow G(s_i, a)$

$R_a + r_i$

for $v \in G$

$V_{a,o_i,v} = V_{a,o_i,v} + \text{Simulate}(G, v, s', L)$

for o in O

$V_{a,o} = \max_{v \in G} V_{a,o,v}$

$v_{a,o} = \text{argmax}_{v \in G} V_{a,o,v}$

$V_a = R_a + \gamma \sum_o V_{ao}/N$

$V^* = \max_a V_a$

$a^* = \text{argmax}\ V_a$

add new node to G labeled with a^*

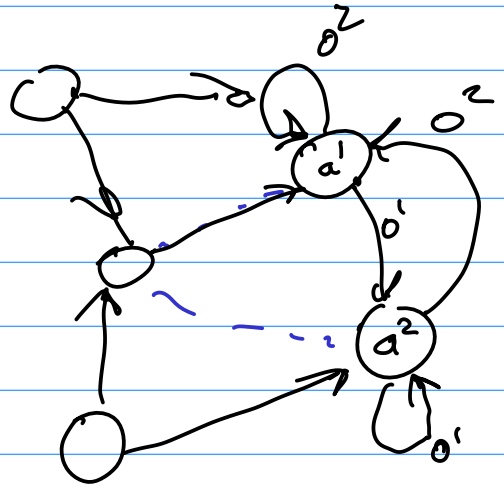

V = value

v = node in graph