

ASEN 5519-003 Decision Making under Uncertainty

Homework 6: POMDPs

March 22, 2021

1 Exercises

Question 1. (35 pts)

Write the following three elements of a POMDP solver:

1. A belief **Updater** and the associated **update** function to calculate the Bayesian belief update¹.
2. A **Policy struct** and associated **action** function that chooses an action based on a list of alpha vectors.
3. A function that calculates the QMDP alpha vectors for a POMDP and returns them as an object of the **Policy** type described above.

The starter code contains the skeleton code for these three items. Use your QMDP code to solve the **TigerPOMDP** from the **POMDPModels.jl** package and evaluate the resulting policy and a near-optimal policy calculated with the **SARSOP.jl** package using Monte Carlo simulations. (The deliverables for this problem are just the code and the average returns from the simulations.)

Question 2. (25 pts)

Evaluate the following three policies on the cancer problem that you created in Homework 5:

1. The QMDP policy obtained using your code from Question 1 above.
2. A heuristic policy *that outperforms the QMDP policy*².
3. A near-optimal policy calculated by the **SARSOP** solver from the **SARSOP.jl** package.

First, report the results of these simulations in a table, then write a short paragraph answering the following question: Both the tiger and cancer problems are similar in that they have information-gathering actions. However, the gap between the QMDP and optimal solutions in one of the problems is much larger. Which problem has the larger gap, and why?

2 Challenge Problem

Question 3. (Lasertag POMDP, 40 pts)

In this problem, you will find a policy and belief updater for the laser tag POMDP model `HW6.LaserTagPOMDP()`. In this POMDP, a robot seeks to tag a moving target in a grid world with obstacles. The state space consists of all positions the robot and the target can take, and the problem ends with a reward of 100 as soon as they occupy the same cell. There are five actions, `:up`, `:down`, `:left`, `:right`, and `:measure`. The `:measure` action has a cost of -2 and gives the robot returns from lasers pointed in the four directions indicating the *exact* distance to the first wall, obstacle, or target that the laser encounters. When any other action is taken,

¹This can be a discrete Bayesian filter or a particle filter

²Hint: you may want to use the QMDP policy within your heuristic policy, i.e. take the QMDP actions some of the time.

the reward is -1 and the laser return for each direction is noisy; it is uniformly distributed between 0 and the distance to the nearest object. The `POMDPs.jl` interface including `POMDPModelTools.render` can be used to further explore the problem.

- a) Submit a `Policy` or a tuple containing a `Policy` and a `Updater` to the `HW6.evaluate` function. A score of 35 will get full credit. You can use or modify your QMDP code from Question 1, or you are encouraged to use any tools available to you - any `POMDPs.jl` solvers, deep reinforcement learning, a modification of your MCTS code from earlier homework, or heuristic policies are all acceptable. The `Policy` object may be a solution that was calculated offline, or an online planner.
- b) Write a short paragraph describing your approach.