

Last Time

Last Time

- Does value iteration always converge?
- Is the value function unique?

Yes!
Because
 B is a
contraction

any number of optimal policies

every optimal policy achieves U^*

$$\begin{array}{l} \pi^1 \\ \pi^2 \end{array} \quad \begin{array}{l} U^{\pi^1} = U^* \\ U^{\pi^2} = U^* \end{array}$$


$$A = \{ \text{left}, \text{right}, \text{right} \}$$

Guiding Questions

Guiding Questions

- What are the differences between *online* and *offline* solutions?
- Are there solution techniques that require computation time *independent* of the state space size?

Curse of Dimensionality



Curse of Dimensionality

1 dimension, 5 segments

$$|\mathcal{S}| = 5$$




Curse of Dimensionality

1 dimension, 5 segments

$$|\mathcal{S}| = 5$$

2 dimensions, 5 segments

$$|\mathcal{S}| = 25$$



Curse of Dimensionality

1 dimension, 5 segments


$$|\mathcal{S}| = 5$$

2 dimensions, 5 segments

$$|\mathcal{S}| = 25$$

3 dimensions, 5 segments

$$|\mathcal{S}| = 125$$



Curse of Dimensionality

1 dimension, 5 segments


$$|\mathcal{S}| = 5$$

2 dimensions, 5 segments

$$|\mathcal{S}| = 25$$

3 dimensions, 5 segments

$$|\mathcal{S}| = 125$$



$$n \text{ dimensions, } k \text{ segments} \rightarrow |\mathcal{S}| = k^n$$

Offline vs Online Solutions

Offline

Online

Offline vs Online Solutions

Offline

- Before Execution: find V^*/Q^*

Online

Offline vs Online Solutions

Offline


- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$

Online

Offline vs Online Solutions

Offline

- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




Online

Offline vs Online Solutions

Offline

- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




Online

- Before Execution: <nothing>

Offline vs Online Solutions

Offline

- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)

Offline vs Online Solutions


Offline

- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$



Online


- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)



Offline vs Online Solutions


Offline

- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$



Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)




- Why?

Offline vs Online Solutions


Offline

- Before Execution: find V^*/Q^*
- During Execution: $\pi^*(s) = \operatorname{argmax} Q^*(s, a)$




Online

- Before Execution: <nothing>
- During Execution: Consider actions and their consequences (everything)



- Why?
- Online methods are insensitive to the size of S !

One Step Lookahead




```
randstep( $\mathcal{P}$ ::MDP, s, a) =  $\mathcal{P}$ .TR(s, a)
```

```
function rollout( $\mathcal{P}$ , s,  $\pi$ , d)
    ret = 0.0
    for t in 1:d
        a =  $\pi$ (s)
        s, r = randstep( $\mathcal{P}$ , s, a)
        ret +=  $\mathcal{P}$ . $\gamma$ ^(t-1) * r
    end
    return ret
end
```

```
function ( $\pi$ ::RolloutLookahead)(s)
     $U(s)$  = rollout( $\pi$ . $\mathcal{P}$ , s,  $\pi$ . $\pi$ ,  $\pi$ .d)
    return greedy( $\pi$ . $\mathcal{P}$ , U, s).a
end
```

```
function greedy( $\mathcal{P}$ ::MDP, U, s)
    u, a = findmax(a → lookahead( $\mathcal{P}$ , U, s, a),  $\mathcal{P}$ .A)
    return (a=a, u=u)
end
```

Forward Search



$$Q(s, a)$$

$$U = \max Q(s, a)$$

$$Q = R(s, a) + \gamma E[U(s)]$$

```
function forward_search( $\mathcal{P}$ , s, d, U)
  if  $d \leq 0$ 
    return (a=nothing, u=U(s))
  end
  best = (a=nothing, u=-Inf)
   $U'(s) = \text{forward\_search}(\mathcal{P}, s, d-1, U).u$ 
  for a in  $\mathcal{P}.A$ 
     $\rightarrow u = \text{lookahead}(\mathcal{P}, \underline{U'}, s, a)$ 
    if  $u > \text{best}.u$ 
      best = (a=a, u=u)
    end
  end
  return best
end
```




$$O((\underline{|S|} \times |A|)^d)$$


sad!

Forward Search depth

Forward Search depth



Forward Search depth



Sparse Sampling

```
function sparse_sampling( $\mathcal{P}$ , s, d, m, U)
    if d ≤ 0
        return (a=nothing, u=U(s))
    end
    best = (a=nothing, u=-Inf)
    for a in  $\mathcal{P}.\mathcal{A}$ 
        u = 0.0
        for i in 1:m
            s', r = randstep( $\mathcal{P}$ , s, a)
            a', u' = sparse_sampling( $\mathcal{P}$ , s', d-1, m, U)
            u += (r +  $\mathcal{P}.\gamma*u'$ ) / m
        end
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```


Sparse Sampling

```
function sparse_sampling( $\mathcal{P}$ , s, d, m, U)
    if d ≤ 0
        return (a=nothing, u=U(s))
    end
    best = (a=nothing, u=-Inf)
    for a in  $\mathcal{P}.A$ 
        u = 0.0
        for i in 1:m
            s', r = randstep( $\mathcal{P}$ , s, a)
            a', u' = sparse_sampling( $\mathcal{P}$ , s', d-1, m, U)
            u += (r +  $\mathcal{P}.\gamma*u'$ ) / m
        end
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

$$O((m|A|)^d)$$

Sparse Sampling

```
function sparse_sampling( $\mathcal{P}$ , s, d, m, U)
    if d ≤ 0
        return (a=nothing, u=U(s))
    end
    best = (a=nothing, u=-Inf)
    for a in  $\mathcal{P}.A$ 
        u = 0.0
        for i in 1:m
            s', r = randstep( $\mathcal{P}$ , s, a)
            a', u' = sparse_sampling( $\mathcal{P}$ , s', d-1, m, U)
            u += (r +  $\mathcal{P}.\gamma * u'$ ) / m
        end
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

$$O((m|A|)^d)$$

$$|V^{\text{SS}}(s) - V^*(s)| \leq \epsilon$$

Sparse Sampling


```
function sparse_sampling( $\mathcal{P}$ , s, d, m, U)
    if d ≤ 0
        return (a=nothing, u=U(s))
    end
    best = (a=nothing, u=-Inf)
    for a in  $\mathcal{P}.A$ 
        u = 0.0
        for i in 1:m
            s', r = randstep( $\mathcal{P}$ , s, a)
            a', u' = sparse_sampling( $\mathcal{P}$ , s', d-1, m, U)
            u += (r +  $\mathcal{P}.\gamma*u'$ ) / m
        end
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

$$O((m|A|)^d)$$

$$|V^{\text{SS}}(s) - V^*(s)| \leq \epsilon$$

m , ϵ , and d related, but independent of $|S|$

Sparse Sampling



$m=3$

```
function sparse_sampling( $\mathcal{P}$ , s, d, m, U)
    if d ≤ 0
        return (a=nothing, u=U(s))
    end
    best = (a=nothing, u=-Inf)
    for a in  $\mathcal{P}.A$ 
        u = 0.0
        for i in 1:m
            s', r = randstep( $\mathcal{P}$ , s, a)
            a', u' = sparse_sampling( $\mathcal{P}$ , s', d-1, m, U)
            u += (r +  $\mathcal{P}.\gamma * u'$ ) / m
        end
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

$$O((m|A|)^d) \quad |V^{\text{SS}}(s) - V^*(s)| \leq \epsilon \quad m, \epsilon, \text{ and } d \text{ related, but independent of } |S|$$

mad!




<https://www.cis.upenn.edu/~mkearns/papers/sparsesampling-journal.pdf>

not on
exam

Break

Draw the trees produced by the following algorithms for a problem with 2 actions and 3 states:

1. One-step lookahead with rollout
2. Forward search ($d=2$)
3. Sparse sampling ($d=2, m=2$)



Branch and Bound

Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$

```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
     $U'(s)$  = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ ,  $U'$ , s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Branch and Bound


Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$

$\underline{V}(s)$ ○

```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
     $U'(s)$  = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ ,  $U'$ , s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Branch and Bound


Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$



```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
  if d ≤ 0
    return (a=nothing, u=Ulo(s))
  end
  U'(s) = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
  best = (a=nothing, u=-Inf)
  for a in sort( $\mathcal{P}.\mathcal{A}$ , by=a→Qhi(s,a), rev=true)
    if Qhi(s, a) < best.u
      return best # safe to prune
    end
    u = lookahead( $\mathcal{P}$ , U', s, a)
    if u > best.u
      best = (a=a, u=u)
    end
  end
  return best
end
```


Branch and Bound


Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$



```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
    U'(s) = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ , U', s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Branch and Bound


Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$



```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
    U'(s) = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ , U', s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Branch and Bound


Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$



```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
    U'(s) = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ , U', s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Branch and Bound


Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$



```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
    U'(s) = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ , U', s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Branch and Bound


Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$



```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
    U'(s) = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ , U', s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Branch and Bound

Assume you have $\underline{V}(s)$ and $\bar{Q}(s, a)$



```
function branch_and_bound( $\mathcal{P}$ , s, d, Ulo, Qhi)
    if d ≤ 0
        return (a=nothing, u=Ulo(s))
    end
    U'(s) = branch_and_bound( $\mathcal{P}$ , s, d-1, Ulo, Qhi).u
    best = (a=nothing, u=-Inf)
    for a in sort( $\mathcal{P}.A$ , by=a→Qhi(s,a), rev=true)
        if Qhi(s, a) < best.u
            return best # safe to prune
        end
        u = lookahead( $\mathcal{P}$ , U', s, a)
        if u > best.u
            best = (a=a, u=u)
        end
    end
    return best
end
```

Monte Carlo Tree Search (MCTS/UCT)

Monte Carlo Tree Search (MCTS/UCT)

Search

Monte Carlo Tree Search (MCTS/UCT)

Search

Expansion

Monte Carlo Tree Search (MCTS/UCT)

Search

Expansion

Rollout

Monte Carlo Tree Search (MCTS/UCT)

Search

Expansion

Rollout

Backup

Monte Carlo Tree Search (MCTS/UCT)

Search



Expansion

Rollout


Backup

Monte Carlo Tree Search (MCTS/UCT)

Search



Expansion




Rollout


Backup

Monte Carlo Tree Search (MCTS/UCT)

Search



Expansion




Rollout


Backup

Monte Carlo Tree Search (MCTS/UCT)

Search



Expansion



Rollout

Backup

more Common

more theoretically justified


$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \leftarrow \text{or } Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus


start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$

Monte Carlo Tree Search (MCTS/UCT)

Search



Expansion



Rollout

Backup


$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad \text{or} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus


start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$

Monte Carlo Tree Search (MCTS/UCT)


Search



Expansion



Rollout



Backup


$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad \text{or} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus


start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$

Monte Carlo Tree Search (MCTS/UCT)


Search



Expansion



Rollout



Backup


$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad \text{or} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus


start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$

Monte Carlo Tree Search (MCTS/UCT)

Search



Expansion



Rollout



Backup


$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad \text{or} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus


start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$

Monte Carlo Tree Search (MCTS/UCT)

Search




Expansion



Rollout



Backup




$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad \text{or} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus


start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$

Monte Carlo Tree Search (MCTS/UCT)

Search




Expansion



Rollout



Backup




$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad \text{or} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus


start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$

Monte Carlo Tree Search (MCTS/UCT)


Search




Expansion



Rollout




Backup



$$Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \quad \text{or} \quad Q(s, a) + c \frac{N(s)^\beta}{\sqrt{N(s, a)}}$$

low $N(s, a)/N(s)$ = high bonus

start with $c = 2(\bar{V} - \underline{V})$, $\beta = 1/4$



Guiding Questions

Guiding Questions

- What are the differences between online and offline solutions?
- Are there solution techniques that are *independent* of the state space size?