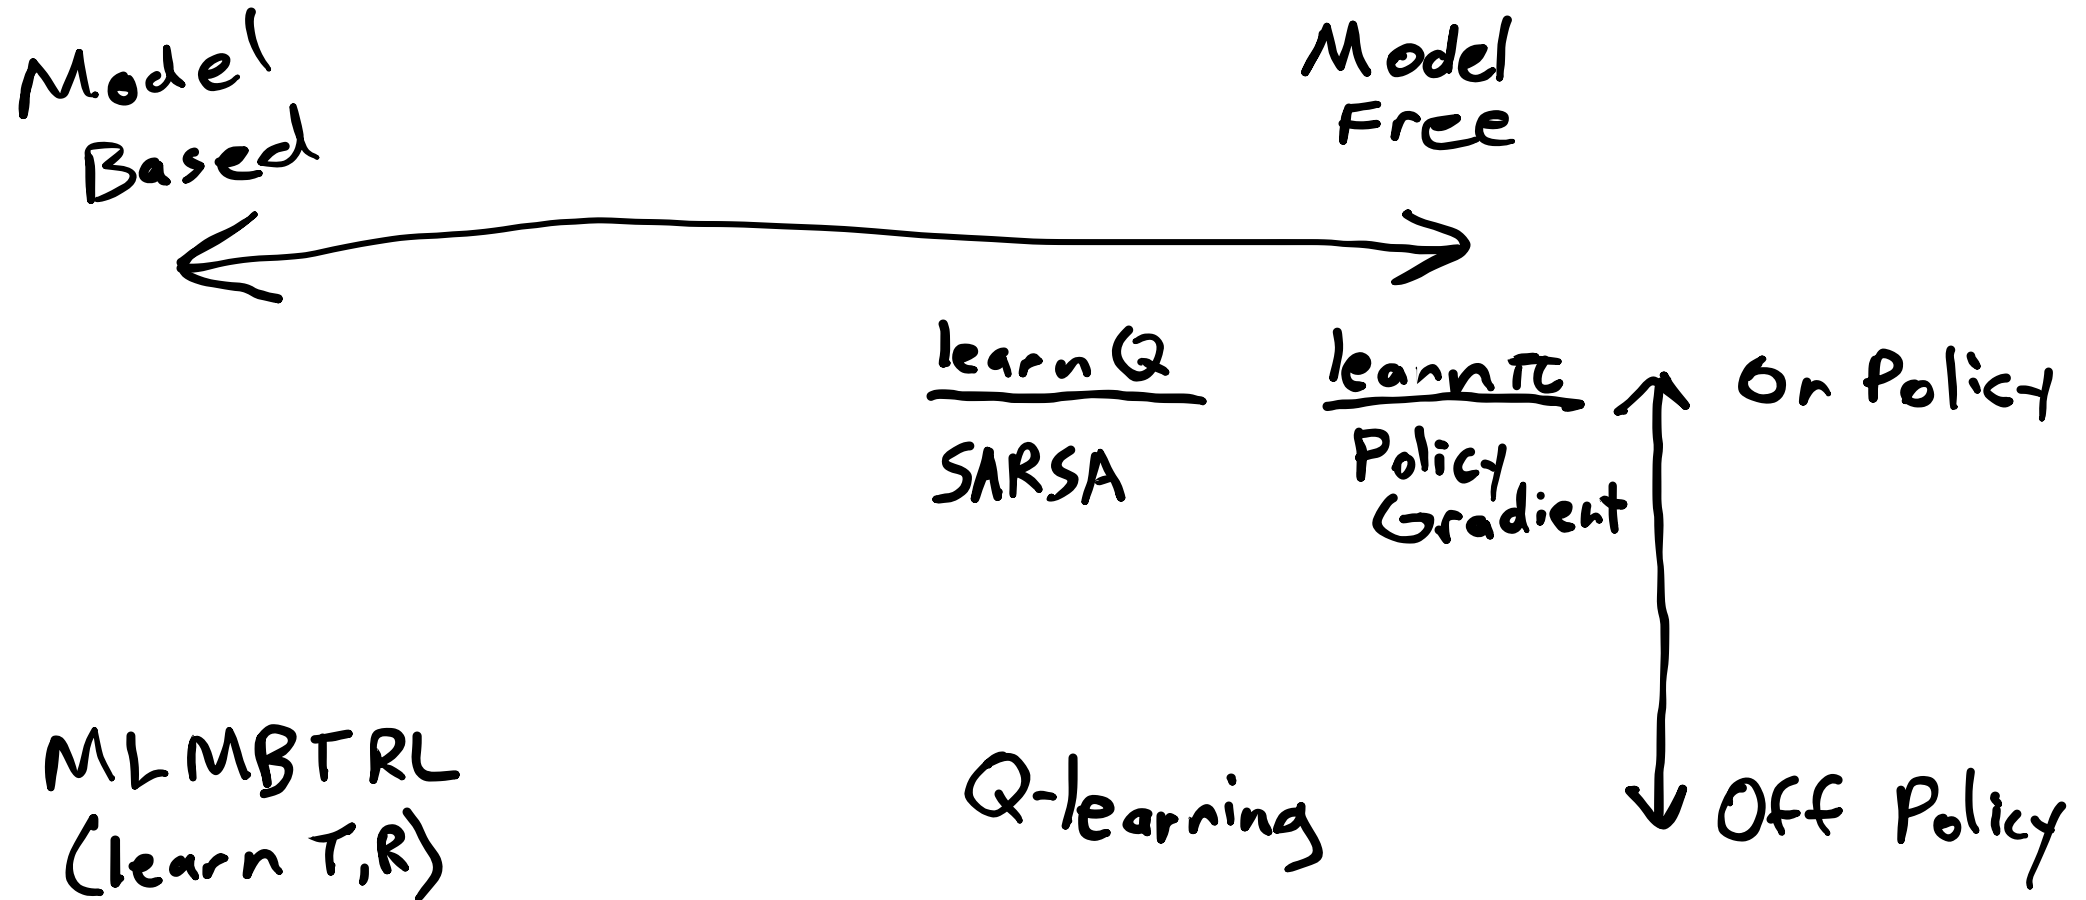


Last Time

- What is Policy Gradient?
- What tricks are needed to make Policy Gradient work?

Map



Today

- Basic On- and Off-Policy **value based** model free RL algorithms
- Tricks for tabular value based RL algorithms
- Understanding of On- vs Off-Policy

Why learn Q?

Incremental Mean Estimation

$$\begin{aligned}\hat{x}_m &= \frac{1}{m} \sum_{i=1}^m x^{(i)} \\ &= \frac{1}{m} \left(x^{(m)} + \sum_{i=1}^{m-1} x^{(i)} \right) \\ &= \frac{1}{m} \left(x^{(m)} + (m-1) \hat{x}_{m-1} \right) \\ &= \hat{x}_{m-1} + \frac{1}{m} \left(x^{(m)} - \hat{x}_{m-1} \right)\end{aligned}$$

```
function simulate!( $\pi$ ::MonteCarloTreeSearch, s, d= $\pi$ .d)
    if d  $\leq$  0
        return  $\pi$ .U(s)
    end
     $\mathcal{P}$ , N, Q, c =  $\pi$ . $\mathcal{P}$ ,  $\pi$ .N,  $\pi$ .Q,  $\pi$ .c
     $\mathcal{A}$ , TR,  $\gamma$  =  $\mathcal{P}$ . $\mathcal{A}$ ,  $\mathcal{P}$ .TR,  $\mathcal{P}$ . $\gamma$ 
    if !haskey(N, (s, first( $\mathcal{A}$ )))
        for a in  $\mathcal{A}$ 
            N[(s,a)] = 0
            Q[(s,a)] = 0.0
        end
    end
    return  $\pi$ .U(s)
end
a = explore( $\pi$ , s)
s', r = TR(s,a)
q = r +  $\gamma$ *simulate!( $\pi$ , s', d-1)
Q[(s,a)] += (q-Q[(s,a)])/N[(s,a)]
end
```

loop

$$\hat{x} \leftarrow \hat{x} + \alpha (x - \hat{x})$$

"Temporal Difference
(TD) Error"

Q Learning

Q learning and SARSA

Q-Learning

$$Q(s, a) \leftarrow 0$$

$$s \leftarrow s_0$$

loop

$$a \leftarrow \operatorname{argmax} Q(s, a) \text{ w.p. } 1 - \epsilon, \quad \operatorname{rand}(A) \text{ o.w.}$$

$$r \leftarrow \operatorname{act!}(\operatorname{env}, a)$$


$$s' \leftarrow \operatorname{observe}(\operatorname{env})$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(\underbrace{r + \gamma \max_{a'} Q(s', a') - Q(s, a)}_{\text{TD}} \right)$$

$$s \leftarrow s'$$

TD

Illustrative Problem




1. After a few episodes, what is $Q(3, a)$ for a in 1-4?
2. After a few episodes, what is $Q(1, L)$?
3. Why is this a problem and what are some possible solutions?

Big Problem: Maximization Bias

Even if all $Q(s', a')$ unbiased, $\max_{a'} Q(s', a')$ is biased!

Solution: Double Q Learning Q_1, Q_2

$$Q_1(s, a) \leftarrow Q_1(s, a) + \alpha \left(r + \gamma Q_2 \left(s', \underset{a'}{\operatorname{argmax}} Q_1(s', a') \right) - Q_1(s, a) \right)$$



Eligibility Traces

Games

Half-Life at 20: why it is the most important shooter ever made

From its opening scenes, Valve's pioneering sci-fi horror game reinvented storytelling and universe building - what made it such a terrifying success?



It taught a whole generation of big-budget game developers how to tell stories' ... the Half-Life box art. Illustration: Valve

SARSA- λ

$$Q(s, a), N(s, a) \leftarrow 0$$

initialize s, a, r, s'

loop

$$a' \leftarrow \operatorname{argmax} Q(s', a) \text{ w.p. } 1 - \epsilon, \quad \text{rand}(A) \text{ o.w.}$$

$$N(s, a) \leftarrow N(s, a) + 1$$

$$\delta \leftarrow r + \gamma Q(s', a') - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \delta N(s, a) \quad \forall s, a$$

$$N(s, a) \leftarrow \gamma \lambda N(s, a)$$

$$s \leftarrow s', \quad a \leftarrow a'$$

$$r \leftarrow \text{act!}(\text{env}, a)$$

$$s' \leftarrow \text{observe}(\text{env})$$

Convergence

- Q learning converges to optimal Q-values w.p. 1
(Sutton and Barto, p. 131)
- SARSA converges to optimal Q-values w.p. 1 ***provided that***
 $\pi \rightarrow \text{greedy}$
(Sutton and Barto, p. 129)

On vs Off-Policy

On Policy

SARSA:

$$Q(s, a) \leftarrow Q(s, a) + \alpha (r + \gamma Q(s', a') - Q(s, a))$$

Off Policy

Q-learning:

$$Q(s, a) \leftarrow Q(s, a) + \alpha (r + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

Will eligibility traces work with Q-learning?

Not easily

Policy Gradient:

$$\theta \leftarrow \theta + \alpha \sum_{k=0}^d \nabla_{\theta} \log \pi_{\theta}(a_k \mid s_k) R(\tau)$$

Today

- Basic On- and Off-Policy **value based** model free RL algorithms
- Tricks for tabular value based RL algorithms
- Understanding of On- vs Off-Policy