

Simple Games

- Games: a mathematical formalism for rational interaction

Simple Games

- Games: a mathematical formalism for rational interaction
- What is the best solution concept? (Nash Equilibrium)

Types of Uncertainty

Types of Uncertainty

Alleatory

Types of Uncertainty

Alleatory



Markov Decision Process

Types of Uncertainty

Alleatory

Epistemic (Static)




Markov Decision Process

Types of Uncertainty

Alleatory

Epistemic (Static)



Markov Decision Process

Reinforcement Learning


Types of Uncertainty

Alleatory



Markov Decision Process

Epistemic (Static)



Reinforcement Learning

Epistemic (Dynamic)



POMDP

Types of Uncertainty

Alleatory



Markov Decision Process

Epistemic (Static)



Reinforcement Learning

Epistemic (Dynamic)



POMDP

Interaction

Types of Uncertainty

Alleatory



Markov Decision Process

Epistemic (Static)



Reinforcement Learning

Epistemic (Dynamic)



POMDP

Interaction



Game

Normal Form Games

Normal Form Games

- Alice and Bob are working on a homework assignment.

Normal Form Games

- Alice and Bob are working on a homework assignment.
- They can either share or withhold their knowledge.

Normal Form Games

- Alice and Bob are working on a homework assignment.
- They can either share or withhold their knowledge.
- If one player shares knowledge, the other benefits greatly, but the sharer also benefits by getting to test their knowledge

Normal Form Games

- Alice and Bob are working on a homework assignment.
- They can either share or withhold their knowledge.
- If one player shares knowledge, the other benefits greatly, but the sharer also benefits by getting to test their knowledge

Alice's Payoffs

		B	
		S	W
A	S	4	2
	W	3	1

Normal Form Games

- Alice and Bob are working on a homework assignment.
- They can either share or withhold their knowledge.
- If one player shares knowledge, the other benefits greatly, but the sharer also benefits by getting to test their knowledge

Alice's Payoffs

		B	
		S	W
A	S	4	2
	W	3	1

Bob's Payoffs

		B	
		S	W
A	S	4	3
	W	2	1

Normal Form Games

- Alice and Bob are working on a homework assignment.
- They can either share or withhold their knowledge.
- If one player shares knowledge, the other benefits greatly, but the sharer also benefits by getting to test their knowledge

Alice's Payoffs

		B	
		S	W
A	S	4	2
	W	3	1

Bob's Payoffs

		B	
		S	W
A	S	4	3
	W	2	1

Bob

Alice

		S	W
Alice	S	4, 4	2, 3
	W	3, 2	1, 1

Normal Form Games

- Alice and Bob are working on a homework assignment.
- They can either share or withhold their knowledge.
- If one player shares knowledge, the other benefits greatly, but the sharer also benefits by getting to test their knowledge

Alice's Payoffs

		B	
		S	W
A	S	4	2
	W	3	1

Bob's Payoffs

		B	
		S	W
A	S	4	3
	W	2	1

Bob

Alice

		S	W
Alice	S	4, 4	2, 3
	W	3, 2	1, 1

Called a **Normal Form, Simple**, or **Bimatrix** Game

Normal Form Games

- Alice and Bob are working on a homework assignment.
- They can either share or withhold their knowledge.
- If one player shares knowledge, the other benefits greatly, but the sharer also benefits by getting to test their knowledge

Alice's Payoffs

		B	
		S	W
A	S	4	2
	W	3	1

Bob's Payoffs

		B	
		S	W
A	S	4	3
	W	2	1

Bob

		Bob	
		S	W
Alice	S	4, 4	2, 3
	W	3, 2	1, 1

Called a **Normal Form, Simple**, or **Bimatrix** Game

Question for today: What **solution concept** should we use for games?

Dominant Strategies

Dominant Strategies

		Bob	
		S	W
Alice	S	3, 3	2, 2
	W	2, 2	1, 1

Dominant Strategies

		Bob	
		S	W
Alice	S	3, 3	2, 2
	W	2, 2	1, 1

Definitions

- Action $a^i \in A^i$
- Joint Action $a = (a^1, \dots, a^k)$
- All Other Actions
 $a^{-i} = (a^1, \dots, a^{i-1}, a^{i+1}, \dots, a^k)$
- Reward $R^i(a)$
- Joint Reward $R(a) = (R^1(a), \dots, R^k(a))$

Dominant Strategies

		Bob	
		S	W
Alice	S	3, 3	2, 2
	W	2, 2	1, 1

Definitions

- Action $a^i \in A^i$
- Joint Action $a = (a^1, \dots, a^k)$
- All Other Actions
 $a^{-i} = (a^1, \dots, a^{i-1}, a^{i+1}, \dots, a^k)$
- Reward $R^i(a)$
- Joint Reward $R(a) = (R^1(a), \dots, R^k(a))$

Deterministic Best Response:

Action a^i is a deterministic best response
to a^{-i} if

$$R^i(a^i, a^{-i}) \geq R^i(a^{i'}, a^{-i}) \quad \forall a^{i'}$$

Dominant Strategies

		Bob	
		S	W
Alice	S	3, 3	2, 2
	W	2, 2	1, 1

- **Dominant (Pure) Strategy:** Action a is a dominant strategy if it is a best response to every action taken by the other player.

Definitions

- Action $a^i \in A^i$
- Joint Action $a = (a^1, \dots, a^k)$
- All Other Actions
 $a^{-i} = (a^1, \dots, a^{i-1}, a^{i+1}, \dots, a^k)$
- Reward $R^i(a)$
- Joint Reward $R(a) = (R^1(a), \dots, R^k(a))$

Deterministic Best Response:

Action a^i is a deterministic best response to a^{-i} if

$$R^i(a^i, a^{-i}) \geq R^i(a^{i'}, a^{-i}) \quad \forall a^{i'}$$

Dominant Strategies

		Bob	
		S	W
Alice	S	3, 3	2, 2
	W	2, 2	1, 1

- **Dominant (Pure) Strategy:** Action a is a dominant strategy if it is a best response to every action taken by the other player.
- **Dominant Strategy Equilibrium:** Every player plays a dominant strategy

Definitions

- Action $a^i \in A^i$
- Joint Action $a = (a^1, \dots, a^k)$
- All Other Actions
 $a^{-i} = (a^1, \dots, a^{i-1}, a^{i+1}, \dots, a^k)$
- Reward $R^i(a)$
- Joint Reward $R(a) = (R^1(a), \dots, R^k(a))$

Deterministic Best Response:

Action a^i is a deterministic best response to a^{-i} if

$$R^i(a^i, a^{-i}) \geq R^i(a^{i'}, a^{-i}) \quad \forall a^{i'}$$

Dominant Strategies

		Bob	
		S	W
Alice	S	3, 3	2, 2
	W	2, 2	1, 1

- **Dominant (Pure) Strategy:** Action a is a dominant strategy if it is a best response to every action taken by the other player.
- **Dominant Strategy Equilibrium:** Every player plays a dominant strategy

Definitions

- Action $a^i \in A^i$
- Joint Action $a = (a^1, \dots, a^k)$
- All Other Actions
 $a^{-i} = (a^1, \dots, a^{i-1}, a^{i+1}, \dots, a^k)$
- Reward $R^i(a)$
- Joint Reward $R(a) = (R^1(a), \dots, R^k(a))$

Deterministic Best Response:

Action a^i is a deterministic best response to a^{-i} if

$$R^i(a^i, a^{-i}) \geq R^i(a^{i'}, a^{-i}) \quad \forall a^{i'}$$

Is the dominant strategy equilibrium always the best outcome for the players?

A more surprising example: The Prisoner's Dilemma

A more surprising example: The Prisoner's Dilemma

- 2 criminals are captured

A more surprising example:

The Prisoner's Dilemma

- 2 criminals are captured
- Each can either keep silent or testify
 - other keeps silent -> minor conviction (1 year)
 - other testifies -> major conviction: 4 years
 - testify -> 1 year removed from sentence

A more surprising example:

The Prisoner's Dilemma

- 2 criminals are captured
- Each can either keep silent or testify
 - other keeps silent -> minor conviction (1 year)
 - other testifies -> major conviction: 4 years
 - testify -> 1 year removed from sentence

Player 1

Player 2

	S	T
S	-1, -1	-4, 0
T	0, -4	-3, -3

A more surprising example:

The Prisoner's Dilemma

- 2 criminals are captured
- Each can either keep silent or testify
 - other keeps silent -> minor conviction (1 year)
 - other testifies -> major conviction: 4 years
 - testify -> 1 year removed from sentence

Player 1

Player 2

	S	T
S	-1, -1	-4, 0
T	0, -4	-3, -3

- Dominant strategy for both players is to testify

A more surprising example:

The Prisoner's Dilemma

- 2 criminals are captured
- Each can either keep silent or testify
 - other keeps silent -> minor conviction (1 year)
 - other testifies -> major conviction: 4 years
 - testify -> 1 year removed from sentence

Player 1

Player 2

	S	T
S	-1, -1	-4, 0
T	0, -4	-3, -3

- Dominant strategy for both players is to testify
- Dominant strategy equilibrium is a very bad social result (for the criminals)

A more surprising example:

The Prisoner's Dilemma

- 2 criminals are captured
- Each can either keep silent or testify
 - other keeps silent -> minor conviction (1 year)
 - other testifies -> major conviction: 4 years
 - testify -> 1 year removed from sentence

Player 1

Player 2

	S	T
S	-1, -1	-4, 0
T	0, -4	-3, -3


- Dominant strategy for both players is to testify
- Dominant strategy equilibrium is a very bad social result (for the criminals)

Do all games have a dominant strategy equilibrium?

Collision Avoidance Game

Collision Avoidance Game

Example: Airborne Collision Avoidance



Player 1

Player 2


Up

Down

	Up	Down
Up	Collision -5, -5	-1, 0
Down	0, -1	Collision -4, -4

Collision Avoidance Game

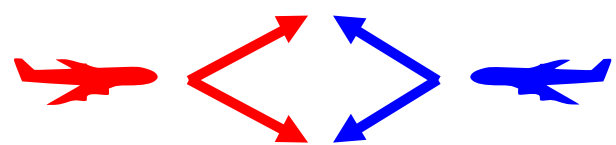
Example: Airborne Collision Avoidance



Pure Nash Equilibrium: All players play a deterministic best response.

Collision Avoidance Game

Example: Airborne Collision Avoidance



Player 2

	Up	Down
Up	Collision -5, -5	-1, 0
Down	0, -1	Collision -4, -4

Player 1

Pure Nash Equilibrium: All players play a deterministic best response.

Which equilibrium is better?

Do all simple games have a pure Nash equilibrium?

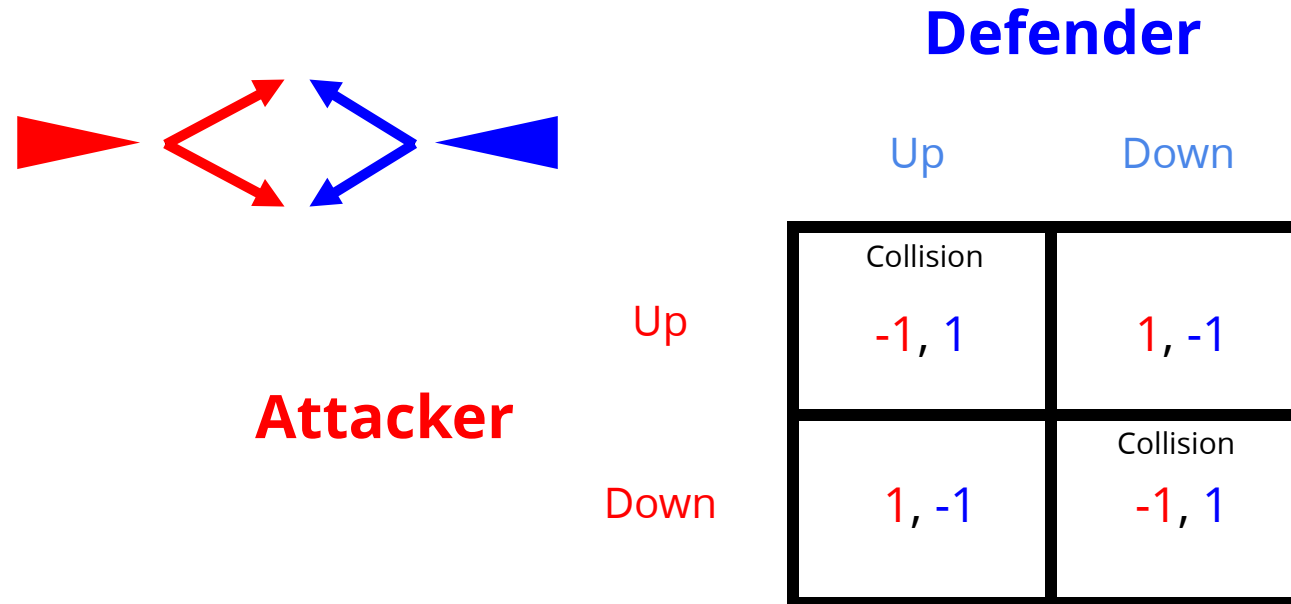
Practice: Find Pure Nash Equilibria

		Player 2		
		a	b	c
Player 1	a	4,4	2,5	0,0
	b	5,2	3,3	0,0
	c	0,0	0,0	10,10

Missile Defense Game

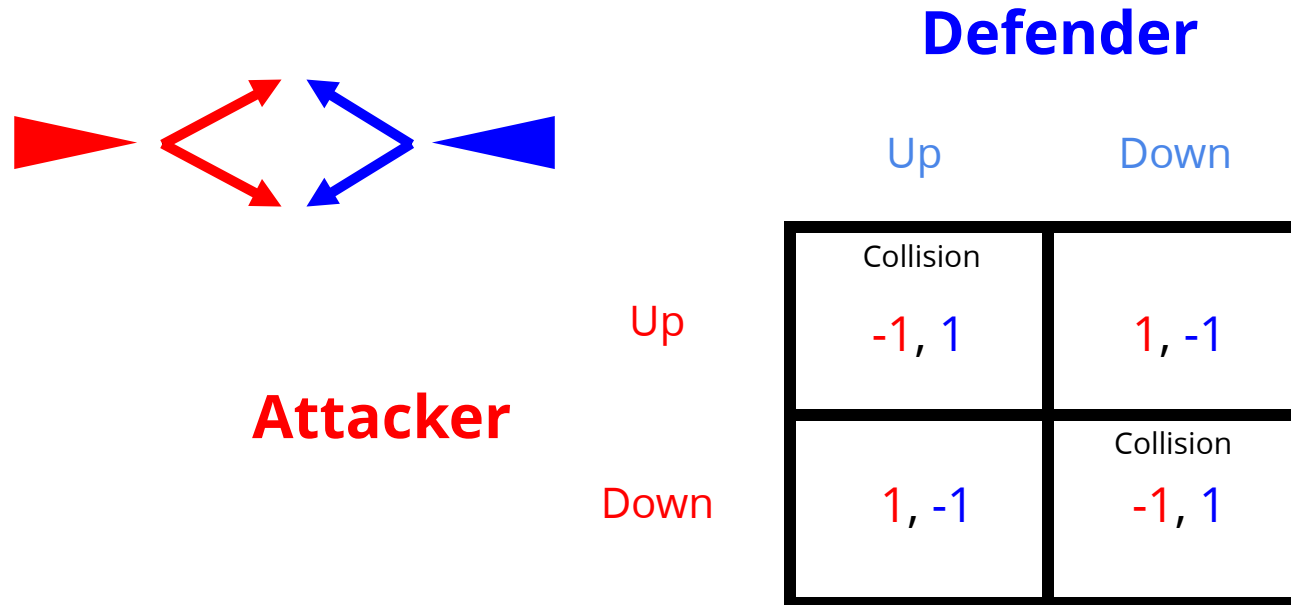
Missile Defense Game

Missile Defense (simplified)



Missile Defense Game


Missile Defense (simplified)



No Pure Nash Equilibrium!

Missile Defense Game

Missile Defense (simplified)



No Pure Nash Equilibrium!

Need a broader solution concept: Mixed Nash equilibrium.

Vocabulary and Notation for Mixed Strategies

Vocabulary and Notation for Mixed Strategies

Single Player

Joint

Vocabulary and Notation for Mixed Strategies

	Single Player	Joint
• Action	$a^i \in A^i$	$a \in A$

Vocabulary and Notation for Mixed Strategies

	Single Player	Joint
• Action	$a^i \in A^i$	$a \in A$
• Policy (strategy)	$\pi^i(a^i)$	$\pi(a) = \prod_i \pi^i(a^i)$

Vocabulary and Notation for Mixed Strategies

	Single Player	Joint
• Action	$a^i \in A^i$	$a \in A$
• Policy (strategy)	$\pi^i(a^i)$	$\pi(a) = \prod_i \pi^i(a^i)$
• Reward	$R^i(a)$	$R(a)$

Vocabulary and Notation for Mixed Strategies

	Single Player	Joint
• Action	$a^i \in A^i$	$a \in A$
• Policy (strategy)	$\pi^i(a^i)$	$\pi(a) = \prod_i \pi^i(a^i)$
• Reward	$R^i(a)$	$R(a)$
• Utility	$U^i(\pi) = \sum_a R^i(a)\pi(a)$	$U(\pi) = \sum_a R(a)\pi(a)$

Vocabulary and Notation for Mixed Strategies

	Single Player	Joint
• Action	$a^i \in A^i$	$a \in A$
• Policy (strategy)	$\pi^i(a^i)$	$\pi(a) = \prod_i \pi^i(a^i)$
• Reward	$R^i(a)$	$R(a)$
• Utility	$U^i(\pi) = \sum_a R^i(a)\pi(a)$	$U(\pi) = \sum_a R(a)\pi(a)$

Two Player Zero Sum:

$$R^1(a) + R^2(a) = 0 \quad \forall a$$

Vocabulary and Notation for Mixed Strategies

	Single Player	Joint
• Action	$a^i \in A^i$	$a \in A$
• Policy (strategy)	$\pi^i(a^i)$	$\pi(a) = \prod_i \pi^i(a^i)$
• Reward	$R^i(a)$	$R(a)$
• Utility	$U^i(\pi) = \sum_a R^i(a)\pi(a)$	$U(\pi) = \sum_a R(a)\pi(a)$

Two Player Zero Sum:

$$R^1(a) + R^2(a) = 0 \quad \forall a$$


Best Response: Given a joint policy of all other agents, π^{-i} , a best response is a policy π^i that satisfies

$$U^i(\pi^i, \pi^{-i}) \geq U^i(\pi^{i'}, \pi^{-i})$$

for all other $\pi^{i'}$.

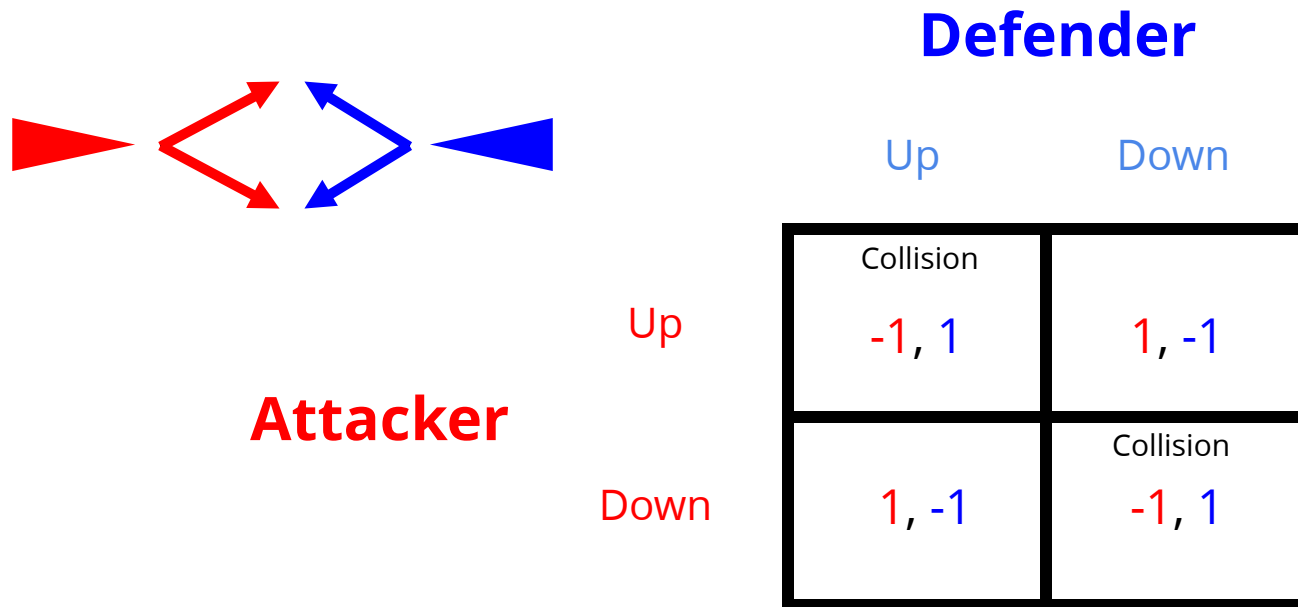
Missile Defense Game

Missile Defense (simplified)



Missile Defense Game

Missile Defense (simplified)



- A ***Nash equilibrium*** is a joint policy in which all agents are following a best response

Rock-paper scissors

1. Guess the Nash Equilibrium argument
2. Make a qualitative argument that this is an NE based on best responses

		agent 2		
		rock	paper	scissors
agent 1	rock	0,0	-1,1	1,-1
	paper	1,-1	0,0	-1,1
	scissors	-1,1	1,-1	0,0

Rock-paper scissors

1. Guess the Nash Equilibrium argument
2. Make a qualitative argument that this is an NE based on best responses

		agent 2		
		rock	paper	scissors
agent 1	rock	0,0	-1,1	1,-1
	paper	1,-1	0,0	-1,1
	scissors	-1,1	1,-1	0,0

Do all simple games have at least one Nash equilibrium?

Rock-paper scissors

1. Guess the Nash Equilibrium argument
2. Make a qualitative argument that this is an NE based on best responses

		agent 2		
		rock	paper	scissors
agent 1	rock	0,0	-1,1	1,-1
	paper	1,-1	0,0	-1,1
	scissors	-1,1	1,-1	0,0

Do all simple games have at least one Nash equilibrium?

Yes!! (might be mixed) 11.2

Every finite game has a Nash Equilibrium

Every finite game has a Nash Equilibrium

EQUILIBRIUM POINTS IN N -PERSON GAMES

BY JOHN F. NASH, JR.*

PRINCETON UNIVERSITY

Communicated by S. Lefschetz, November 16, 1949

One may define a concept of an n -person game in which each player has a finite set of pure strategies and in which a definite set of payments to the n players corresponds to each n -tuple of pure strategies, one strategy being taken for each player. For mixed strategies, which are probability distributions over the pure strategies, the pay-off functions are the expectations of the players, thus becoming polylinear forms in the probabilities with which the various players play their various pure strategies.

Any n -tuple of strategies, one for each player, may be regarded as a point in the product space obtained by multiplying the n strategy spaces of the players. One such n -tuple counters another if the strategy of each player in the countering n -tuple yields the highest obtainable expectation for its player against the $n - 1$ strategies of the other players in the countered n -tuple. A self-countering n -tuple is called an equilibrium point.

The correspondence of each n -tuple with its set of countering n -tuples gives a one-to-many mapping of the product space into itself. From the definition of countering we see that the set of countering points of a point is convex. By using the continuity of the pay-off functions we see that the graph of the mapping is closed. The closedness is equivalent to saying: if P_1, P_2, \dots and $Q_1, Q_2, \dots, Q_n, \dots$ are sequences of points in the product space where $Q_n \rightarrow Q$, $P_n \rightarrow P$ and Q_n counters P_n then Q counters P .

Since the graph is closed and since the image of each point under the mapping is convex, we infer from Kakutani's theorem¹ that the mapping has a fixed point (i.e., point contained in its image). Hence there is an equilibrium point.

In the two-person zero-sum case the "main theorem"² and the existence of an equilibrium point are equivalent. In this case any two equilibrium points lead to the same expectations for the players, but this need not occur in general.

* The author is indebted to Dr. David Gale for suggesting the use of Kakutani's theorem to simplify the proof and to the A. E. C. for financial support.

¹ Kakutani, S., *Duke Math. J.*, 8, 457-459 (1941).

² Von Neumann, J., and Morgenstern, O., *The Theory of Games and Economic Behaviour*, Chap. 3, Princeton University Press, Princeton, 1947.

Every finite game has a Nash Equilibrium

Kakutani's fixed-point theorem

A correspondence $f: X \rightarrow X$ has a fixed point (i.e., $\mathbf{x} \in f(\mathbf{x})$ for some $\mathbf{x} \in X$) if all of the following conditions hold.

- (1) X is a non-empty, closed, bounded, and convex set.
- (2) $f(\mathbf{x})$ is non-empty for any \mathbf{x} .
- (3) $f(\mathbf{x})$ is convex for any \mathbf{x} .
- (4) The set $\{ (\mathbf{x}, \mathbf{y}) \mid \mathbf{y} \in f(\mathbf{x}) \}$ is closed.

Every finite game has a Nash Equilibrium

Kakutani's fixed-point theorem

A correspondence $f: X \rightarrow X$ has a fixed point (i.e., $\mathbf{x} \in f(\mathbf{x})$ for some $\mathbf{x} \in X$) if all of the following conditions hold.

- (1) X is a non-empty, closed, bounded, and convex set.
- (2) $f(\mathbf{x})$ is non-empty for any \mathbf{x} .
- (3) $f(\mathbf{x})$ is convex for any \mathbf{x} .
- (4) The set $\{ (\mathbf{x}, \mathbf{y}) \mid \mathbf{y} \in f(\mathbf{x}) \}$ is closed.

- Let x be a strategy profile, π .

Every finite game has a Nash Equilibrium

Kakutani's fixed-point theorem

A correspondence $f: X \rightarrow X$ has a fixed point (i.e., $\mathbf{x} \in f(\mathbf{x})$ for some $\mathbf{x} \in X$) if all of the following conditions hold.

- (1) X is a non-empty, closed, bounded, and convex set.
- (2) $f(\mathbf{x})$ is non-empty for any \mathbf{x} .
- (3) $f(\mathbf{x})$ is convex for any \mathbf{x} .
- (4) The set $\{ (\mathbf{x}, \mathbf{y}) \mid \mathbf{y} \in f(\mathbf{x}) \}$ is closed.

- Let x be a strategy profile, π .
- Let f be BR , that is, the best response operator

Every finite game has a Nash Equilibrium

Kakutani's fixed-point theorem

A correspondence $f: X \rightarrow X$ has a fixed point (i.e., $\mathbf{x} \in f(\mathbf{x})$ for some $\mathbf{x} \in X$) if all of the following conditions hold.

- (1) X is a non-empty, closed, bounded, and convex set.
- (2) $f(\mathbf{x})$ is non-empty for any \mathbf{x} .
- (3) $f(\mathbf{x})$ is convex for any \mathbf{x} .
- (4) The set $\{ (\mathbf{x}, \mathbf{y}) \mid \mathbf{y} \in f(\mathbf{x}) \}$ is closed.

- Let x be a strategy profile, π .
- Let f be BR , that is, the best response operator
- A fixed point of BR is a Nash Equilibrium

Every finite game has a Nash Equilibrium

Kakutani's fixed-point theorem

A correspondence $f: X \rightarrow X$ has a fixed point (i.e., $\mathbf{x} \in f(\mathbf{x})$ for some $\mathbf{x} \in X$) if all of the following conditions hold.

- (1) X is a non-empty, closed, bounded, and convex set.
- (2) $f(\mathbf{x})$ is non-empty for any \mathbf{x} .
- (3) $f(\mathbf{x})$ is convex for any \mathbf{x} .
- (4) The set $\{ (\mathbf{x}, \mathbf{y}) \mid \mathbf{y} \in f(\mathbf{x}) \}$ is closed.

- Let x be a strategy profile, π .
- Let f be BR , that is, the best response operator
- A fixed point of BR is a Nash Equilibrium
- The BR operator and policy space for finite games meet the conditions above

Every finite game has a Nash Equilibrium

Kakutani's fixed-point theorem

A correspondence $f: X \rightarrow X$ has a fixed point (i.e., $\mathbf{x} \in f(\mathbf{x})$ for some $\mathbf{x} \in X$) if all of the following conditions hold.

- (1) X is a non-empty, closed, bounded, and convex set.
- (2) $f(\mathbf{x})$ is non-empty for any \mathbf{x} .
- (3) $f(\mathbf{x})$ is convex for any \mathbf{x} .
- (4) The set $\{ (\mathbf{x}, \mathbf{y}) \mid \mathbf{y} \in f(\mathbf{x}) \}$ is closed.

- Let x be a strategy profile, π .
- Let f be BR , that is, the best response operator
- A fixed point of BR is a Nash Equilibrium
- The BR operator and policy space for finite games meet the conditions above
- BR has a fixed point for every finite game, i.e. every finite game has a Nash Equilibrium

General approach to find Nash Equilibria


General approach to find Nash Equilibria

$$\begin{array}{ll}\text{minimize}_{\pi, U} & \sum_i (U^i - U^i(\pi)) \\ \text{subject to} & U^i \geq U^i(a^i, \pi^{-i}) \text{ for all } i, a^i \\ & \sum_{a^i} \pi^i(a^i) = 1 \text{ for all } i \\ & \pi^i(a^i) \geq 0 \text{ for all } i, a^i\end{array}$$

General approach to find Nash Equilibria

$$\begin{aligned}
 &\underset{\pi, U}{\text{minimize}} && \sum_i \left(U^i - U^i(\pi) \right) \\
 &\text{subject to} && U^i \geq U^i(a^i, \pi^{-i}) \text{ for all } i, a^i \\
 &&& \sum_{a^i} \pi^i(a^i) = 1 \text{ for all } i \\
 &&& \pi^i(a^i) \geq 0 \text{ for all } i, a^i
 \end{aligned}$$

Topology of bimatrix games:



Algorithms that use best response

Iterated Best Response: randomly cycle between agents who play the best response for the current policy (converges to Nash for certain narrow classes of games)

Fictitious Play:

1. Estimate maximum likelihood policies for opponents:

$$\pi^j(a^j) \propto N(j, a^j)$$

2. Play best response to estimated policy

(converges to Nash for wider class of games, notably zero-sum)

Battle of the Sexes

Battle of the Sexes

- Gabby and Max are going on a date

Battle of the Sexes

- Gabby and Max are going on a date
- Gabby wants to go to a football game

Battle of the Sexes

- Gabby and Max are going on a date
- Gabby wants to go to a football game
- Max wants to go to a movie (He is a rom-com superfan)

Battle of the Sexes

- Gabby and Max are going on a date
- Gabby wants to go to a football game
- Max wants to go to a movie (He is a rom-com superfan)

A hand-drawn blue payoff matrix for the Battle of the Sexes game. The matrix is a square divided into four quadrants by a horizontal and vertical line. The labels 'G' and 'M' are written in blue ink around the matrix. 'G' is written to the left of the matrix, and 'M' is written above the matrix. The quadrants are labeled as follows: top-left is 'G', top-right is 'M', bottom-left is 'G', and bottom-right is 'M'.

	G	M
G	G	M
M	G	M

Battle of the Sexes

- Gabby and Max are going on a date
- Gabby wants to go to a football game
- Max wants to go to a movie (He is a rom-com superfan)

		M	
		G	M
G	G	2, 1	0, 0
	M	0, 2	1, 2

Battle of the Sexes

- Gabby and Max are going on a date
- Gabby wants to go to a football game
- Max wants to go to a movie (He is a rom-com superfan)

Correlated Equilibrium

- A *correlated joint policy* is a single distribution over the joint actions of all agents.
- A *correlated equilibrium* is a correlated joint policy where no agent i can increase their expected utility by deviating from their current action to another.

A hand-drawn payoff matrix for the Battle of the Sexes game. The matrix is a 2x2 grid with 'G' and 'M' as row and column headers. The payoffs are written in the cells: (G, G) is 2,1; (G, M) is 0,0; (M, G) is 0,2; and (M, M) is 1,2.

	G	M
G	2,1	0,0
M	0,2	1,2

Battle of the Sexes

- Gabby and Max are going on a date
- Gabby wants to go to a football game
- Max wants to go to a movie (He is a rom-com superfan)

Correlated Equilibrium

- A *correlated joint policy* is a single distribution over the joint actions of all agents.
- A *correlated equilibrium* is a correlated joint policy where no agent i can increase their expected utility by deviating from their current action to another.

A hand-drawn payoff matrix for the Battle of the Sexes game. The matrix is a square divided into four quadrants. The columns are labeled 'G' and 'M' at the top, and the rows are labeled 'G' and 'M' on the left. The payoffs are written in the quadrants: (G, G) is 2,1; (G, M) is 0,0; (M, G) is 0,2; and (M, M) is 1,2.

	G	M
G	2,1	0,0
M	0,2	1,2

- Easier to find than Nash equilibrium (Linear Program)_{5.7}

Recap

Recap

- Games provide a mathematical framework for analyzing interaction between rational agents

Recap

- Games provide a mathematical framework for analyzing interaction between rational agents
- Games may not have a single "optimal" solution; instead there are equilibria

Recap

- Games provide a mathematical framework for analyzing interaction between rational agents
- Games may not have a single "optimal" solution; instead there are equilibria
- If every player is playing a best response, that joint policy is a Nash Equilibrium

Recap

- Games provide a mathematical framework for analyzing interaction between rational agents
- Games may not have a single "optimal" solution; instead there are equilibria
- If every player is playing a best response, that joint policy is a Nash Equilibrium
- Every finite game has at least one Nash Equilibrium (pure or mixed)