POMDP Model Approx. "Formulation" Approx

Last two Lectures: Numerically Approximate Solutions
- Offline PBVI / α vectors
- Online tree search

DESPOT α

Model Approximation
- Solve a slightly different opt. prob. "model"
- Toolbox
- when to use

DESPOT-α

~POMCPOW: MCTS + Prog. Wid. + Weighted Particle Beliefs

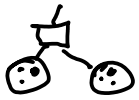

DESPOT: Heuristic Search + Scenarios

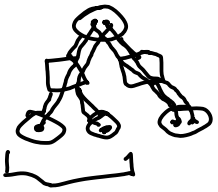

DESPOT-α

All scenarios in *all* belief node weighted by Z

$s', o, r = G(s, a, \Phi_{i,t})$

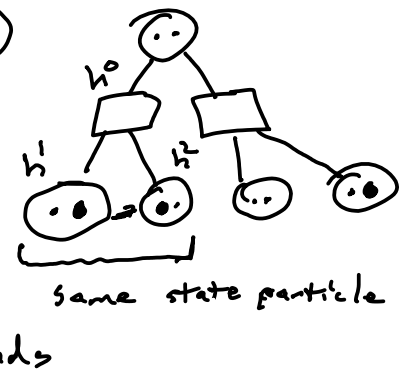

Each Scenario i

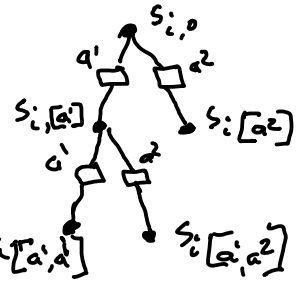

$$U(h') = \sum_{i=1}^{K} w_i^{h'} U(s_{i,h^0})$$

$$L(h') = \sum_{i=1}^{K} w_i^{h'} \cdot \alpha^{p}[s_{i,h^0}]$$

any suboptimal plan

Model Approximations

QMDP

α-vectors $\sim$ Q values $\quad \alpha^p[s] = Q^p(s, p[s])$

$$\alpha^a[s] = Q_{MDP}(s,a)$$

↑ solution to underlying MDP

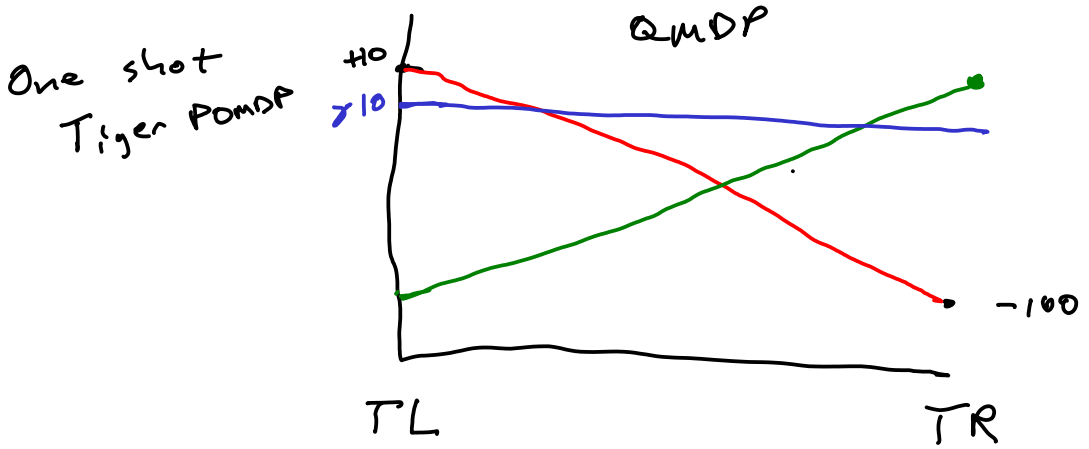

POMDP objective

$$\pi^* = \underset{\pi: B \to A}{\arg\max}\, E\left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t))\right]$$

QMDP

$$\pi_{QMDP}(b) = \underset{a \in A}{\arg\max} E_{s\sim b}[Q_{MDP}(s,a)]$$

- belief at current time
  - ↳ F.O. state in future
- Optimistic

ACAS X

Fully observable — partially

$s = (\Delta h, \dot{h}_0, \dot{h}_i, adv, ps)$

millions

↑ pilot state $\in$ {responsive, ignoring}

QMDP

$$E_{s\sim b}[Q_{MDP}(s,a)]$$

Name	Desc.	Properties	Usefulness
QMDP	Full obs. after 1 step hindsight kn. of epistemic	Upper Bound on true Value Function	★★★★★
FIB	takes 1 step observation into account	Tighter upper Bound than QMDP	★
Hindsight Optimization	Hindsight knowledge of state + outcome (aleatory + epistemic) uncertainty	Looser upper bound than QMDP	★★★★
Certainty Equivalence	Control as if mean (or median or mode) is true state	Optimal for LQR	★★★★★
Open Loop action sequence - - - - - - no observations	Choose action sequence that optimizes objective in expectation	Good when hard to reduce epistemic	★★★
Last k observations "k-markov"	Pretend last k observations are state and solve MDP	Great for Atari!	★★★★
Most Likely Observation	Plan assuming $b' = \tau(b, a, \hat{o}(b))$	No observation branching Good when Z unimodal	★★★
Epistemic → Aleatory	Assume that partially observable part of s takes a random value at each timestep	Conservative	★★

FIB

$$\pi_{FIB}(b) = \underset{a \in A}{\operatorname{argmax}} \alpha_{FIB}^{a\,T} b$$

$$\alpha_a^{(k+1)}[s] = R(s,a) + \gamma \sum_{o} \max_{a'} \sum_{s'} Z(o|a,s') T(s'|s,a) \alpha_a^{(k)}[s']$$

Hindsight

POMDP objective

$$\pi^* = \underset{\pi: B \to A}{\operatorname{argmax}} E\left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

QMDP

$$\pi_{QMDP}(b) = \underset{a \in A}{\operatorname{argmax}} E_{s \sim b}\left[Q_{MDP}(s,a) \right]$$

HOP

$$V_{hs}(b) = E\left[\max_{(a_1 \ldots a_T)} \sum_{t=0}^{T} \gamma^t R(s_t, a_t) \right]$$

$$\pi_{hs}(b) = \underset{a \in A}{\operatorname{argmax}} E_{s \sim b}\left[R(s_0, a) + \max_{\substack{a_t \\ t \in 1..T}} \sum_{t=1}^{T} \gamma^t R(s_t, a_t) \right]$$

(the second term: in VI)

CE

$$\pi_s : S \to A$$

$$\pi_{CE}(b) = \pi_s\left(E_{s \sim b}[s] \right)$$

O.L.

$$\max_{(a_1 \ldots a_T)} E\left[\sum_{t=0}^{T} \gamma^t R(s_t, a_t) \right]$$

no observation

$$b' = \tau(b, a)$$