Last Time

Actor-Critic
Exploration: RND

This Time

POMDP
Bayesian Filters

How to choose an RL algorithm (According Sergey Levine)

Sample Efficiency Ease of Use/Stability

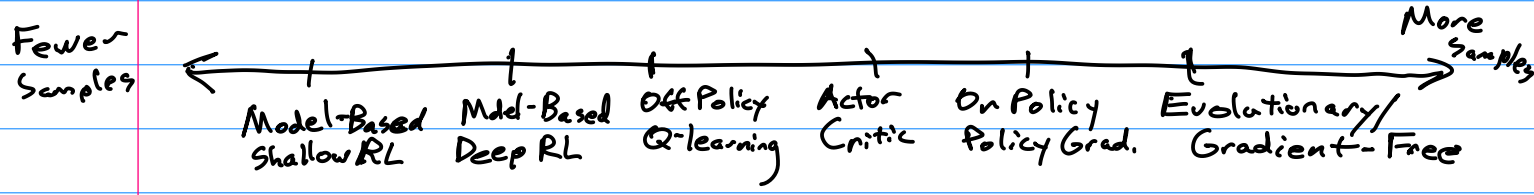

With fast simulator, wall clock time is roughly reversed

POMDPs

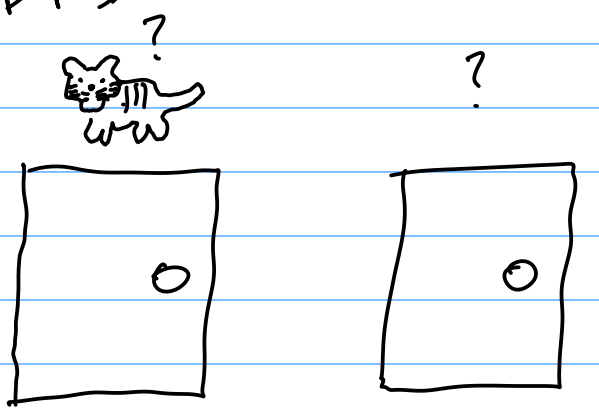

S, A, T, R, O, Z

$S = \{TL, TR\}$

$A = \{OL, OR, L\}$

$T(s'|s,a) = \begin{cases} 1 & \text{if } s' = s \\ 0 & \text{o.w.} \end{cases}$

Reward: +10 if open empty door, -100 if open tiger door

$\gamma = 0.99$

$O = \{TL, TR\}$

$P(o|s) = \begin{cases} 0.85 & \text{if } o = s \\ 0.15 & \text{o.w.} \end{cases} = Z(o|s)$

$\hat{\pi}(TL) = OR$

$\hat{\pi}(TR) = OL$

$V^{\hat{\pi}}(b_0) = 0.85 \cdot 10 + 0.15(-100)$
$= -6.5$

Belief Updating

Hidden Markov Model — HMM

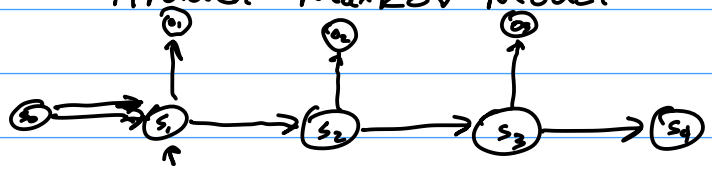

$$P(s_1 \mid o_1) = \frac{P(o_1 \mid s_1) P(s_1)}{P(o_1)} = \frac{P(o_1 \mid s_1) \sum_{s \in S} P(s_1 \mid s_0 = s) P(s_0 = s)}{P(o_1)}$$

$$\propto P(o_1 \mid s_1) \sum_{s \in S} P(s_1 \mid s_0 = s) P(s_0 = s)$$

$b_t(s) = P(s_t = s \mid h_t)$

$h_t = (o_1, \ldots, o_t)$

$b_0(s)$

$$P(s_k \mid h_k) = \frac{P(o_k \mid s_k, h_{k-1}) P(s_k, h_{k-1})}{P(h_k)}$$

$$\propto P(o_k \mid s_k) \sum_{s_{k-1}} P(s_k \mid s_{k-1}, h_{k-1}) P(s_{k-1} \mid h_{k-1}) P(h_{k-1})$$

$$\propto P(o_k \mid s_k) \sum_{s_{k-1}} P(s_k \mid s_{k-1}) \frac{P(s_{k-1} \mid h_{k-1})}{b_{k-1}}$$

$$b'(s) \propto Z(o \mid s') \sum T(s' \mid s) b(s)$$

Belief Update

$O(|S|^2)$

b_0

loop

receive o

for $s' \in S$

$b'(s') \leftarrow Z(o \mid s') \sum_s T(s' \mid s) b(s)$

$b' = b' / \sum_{s'} b'(s')$

$b \leftarrow b'$

$b_0(TL) = 0.5$

$o_1 = TL$ ✓ → $o_1 = TR$

<belief update>

$b_1(TL) = 0.85$ $\quad$ $b_1(TL) = 0.15$

$o_2 = TL$

$b_2(TL) = 0.97$

$o_3 = TR$

$b_3(TL) = 0.85$