

Last Time

Sampling and Inference in Bayes Nets

↑
Input: BN (G, θ)
Output: samples for every variable

Input: BN, values of some
Output: Distributions of others

This Time

Utility Theory

Simple Games

Preferences

$$A \succ B$$

$$A \sim B$$

$$A \succeq B$$

Lottery $[s_1:p_1, s_2:p_2, \dots, s_n:p_n]$

Assumptions

Completeness: one of $A \succ B$, $B \succ A$, or $A \sim B$ holds

Transitivity : if $A \succeq B$ and $B \succeq C$ then $A \succeq C$

Continuity : if $A \succeq C \succeq B \quad \exists p \in (0,1)$ s.t.
 $[A:p, B:1-p] \sim C$

Independence : if $A \succ B$ then for any C and p
 $[A:p, C:1-p] \succeq [B:p, C:1-p]$

If assumptions

$\exists U$ s.t. $U(A) > U(B)$ iff $A \succ B$

$U(A) = U(B)$ iff $A \sim B$

Preference Elicitation

$$U(\underline{s}) = 0 \quad U(\bar{s}) = 1$$

find p such that $[\bar{s}:p, \underline{s}:1-p] \sim S$

$$U(S) = p$$

Pitfalls

Humans are often not rational

- Cognitive Bias
- Certainty Effect
- Framing Effect
- Quantitative measures (\$) don't directly work

\$100 1
\$50 11

Indifferent Risk Neutral

Prefer \$100 50% Risk Seeking

Prefer \$50 100% Risk Averse

Simple Games

"Normal Form" "Matrix"

$\bar{a} \in A$

Joint Action Space


$$A = A^1 \times A^2 \times \dots \times A^k$$

$$I = \{1, \dots, k\}$$

Joint Reward Function

$$R(\bar{a}) = (R^1(\bar{a}), R^2(\bar{a}), \dots, R^k(\bar{a}))$$

		climb	descend
1	climb	$(-5, -5)$	$(-1, 0)$
	descend	$(0, -1)$	$(-4, -4)$



Joint policy $\pi^i(a)$
strategy profile

probability that agent i
plays action a

Pure Strategy if $\pi^i(a) \in \{0, 1\}$

Mixed Strategy otherwise

Breakout Room

What is the best policy?

Group 2

1. Descend

2. Ascend

correlated

Group 1: Communication

1. Fast always ascend

Nash Equilibrium

Response

$-i$ shorthand for $(1, \dots, i-1, i+1, \dots, k)$ $I \setminus i$

$$\bar{a} = (a^i, a^{-i})$$

$$R(\bar{a}) = R(a^i, a^{-i})$$

$$\bar{\pi} = (\pi^i, \pi^{-i})$$

Best Response

$$\pi^i \text{ s.t. } U^i(\pi^i, \pi^{-i}) \geq U^i(\pi^{i'}, \pi^{-i}) \quad \forall \pi^{i'}$$

Nash Equilibrium

A N.E. is a joint policy π in which
all agents follow a best response

OR

A N.E. is a joint policy in which no agent has an incentive to unilaterally switch their policy.

$(-5, -5)$	$(-1, 0)$
$(0, -1)$	$(-4, 4)$

$$\pi^1(\text{climb}) = 1$$

$$\pi^2(\text{climb}) = 0$$

$$\pi^1(\text{climb}) = 0$$

$$\pi^2(\text{climb}) = 1$$

All finite-action₂ games have at least 1 N.E.

R.P.S.

		R	P	S
R		(0,0)	(-1,1)	(1,-1)
P		(1,-1)	(0,0)	(-1,1)
S		(-1,1)	(1,-1)	(0,0)

Uniformly Random

$$\pi^i(R) = \pi^i(P) = \pi^i(S) = \frac{1}{3}$$

		defender	
		climb	descend
attacker	climb	(3, -5)	(-1, 0)
	descend	(0, -1)	(4, -4)

$$\pi^* = \begin{cases} \pi^d(\text{climb}) = 0.5 \\ \pi^a(\text{descend}) = 1.0 \end{cases}$$

finding NE is PPA D-complete : No known polynomial-time solutions

$$\begin{aligned} & \underset{\pi, U}{\text{minimize}} \quad \sum_i (U^i - U^i(\pi)) && 0 \text{ if NE} \\ & \text{subject to} \quad U^i \geq U^i(a^i, \pi^{-i}) && \forall i, a^i \\ & \quad \sum_{a^i} \pi^i(a^i) = 1 && \forall i \\ & \quad \pi^i(a^i) \geq 0 && \forall i, a^i \end{aligned}$$

Correlated Equilibrium

$$\text{For NE. } \pi(\bar{a}) = \prod \pi^i(a^i)$$

$$\sum_{a^i} R^i(a^i, \bar{a}^{-i}) \pi(a^i, \bar{a}^{-i}) \geq \sum_{a^{i'}} R^i(a^{i'}, \bar{a}^{-i}) \pi(a^{i'}, \bar{a}^{-i})$$

Finding a correlated equilibrium involves solving an L.P.

Every NE. is a C.E., but not every C.E. is a N.E.

Things that might converge

- Iterated best response: randomly choose 1 player
solve for best response
cycle through players

- Fictitious Play: Use M.L. estimates of other agents policies, use to choose policy that you play

Humans

Level k response

Level-0 players : random policy

Level- k players calculate best response to level $k-1$ policy

Combine with softmax

\$50 1
\$2 11
,