




POMDPs

- We've been living a lie:



?



`s = observe(env)`

$A = \{OL, OR, L\}$

85% accurate

-1 Listen
+10 opening ~~good~~ ^{safe} door
-100 opening tiger door

(Listen until expected value of one of the doors is positive

(Use Bayes rule to update tiger probabilities based on observations we get from listening

Types of Uncertainty

Types of Uncertainty

Alleatory

Types of Uncertainty

Alleatory



Types of Uncertainty

Alleatory




Epistemic (Static)

Types of Uncertainty

Alleatory



Epistemic (Static)




Types of Uncertainty

Alleatory



Epistemic (Static)




Epistemic (Dynamic)

Types of Uncertainty


Alleatory



Epistemic (Static)



Epistemic (Dynamic)




Types of Uncertainty


Alleatory



Epistemic (Static)



Epistemic (Dynamic)




Interaction

Types of Uncertainty


Alleatory



Epistemic (Static)



Epistemic (Dynamic)




Interaction




Markov Decision Process (MDP)

- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution



Markov Decision Process (MDP)


- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space



Markov Decision Process (MDP)


$$(S, A, T, R, \gamma)$$

- S - State space
- $T : S \times \mathcal{A} \times S \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : S \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward



Markov Decision Process (MDP)

- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward



Alleatory

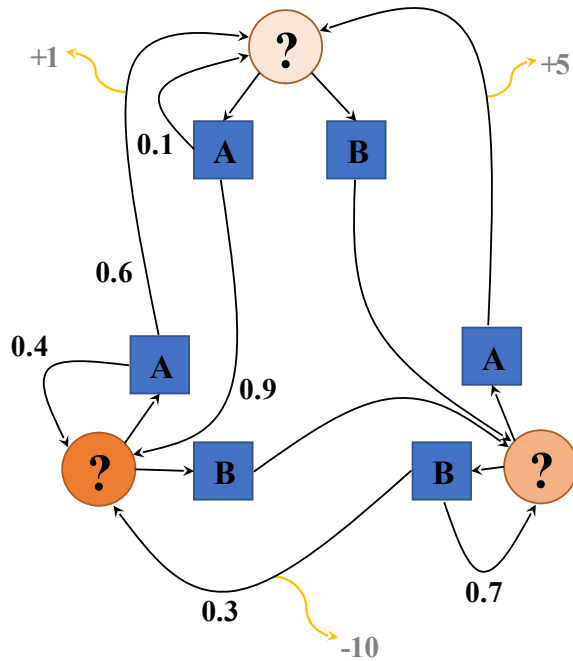
RL

Static Epistemic

R, T


Partially Observable Markov Decision Process (POMDP)

- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward




Partially Observable Markov Decision Process (POMDP)

- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward
- \mathcal{O} - Observation space



Partially Observable Markov Decision Process (POMDP)



- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward
- \mathcal{O} - Observation space
- $Z : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \times \mathcal{O} \rightarrow \mathbb{R}$ - Observation probability distribution


$$o \in \mathcal{O}$$

$$\mathcal{O} = \{TL, TR\}$$

$$Z(o | s, a, s')$$

$$Z(o | a, s')$$


Partially Observable Markov Decision Process (POMDP)



Alleatory

- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward
- \mathcal{O} - Observation space
- $Z : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \times \mathcal{O} \rightarrow \mathbb{R}$ - Observation probability distribution

Partially Observable Markov Decision Process (POMDP)



- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward
- \mathcal{O} - Observation space
- $Z : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \times \mathcal{O} \rightarrow \mathbb{R}$ - Observation probability distribution

Alleatory

Epistemic (Static)

Epistemic (Dynamic)

Tiger POMDP Definition

$$S = \{TL, TR\}$$

$$A = \{OL, OR, L\}$$

$$O = \{TL, TR\}$$

$$R(s,a) = \begin{cases} -1 & \text{if } a=L \\ +10 & \text{if } a \in \{OL, OR\} \text{ and } a \neq \overline{s} \\ -100 & \text{if } a \in \{OL, OR\} \text{ and } a = \overline{s} \end{cases}$$

$$T(s'|s,a) = \begin{cases} 1 & \text{if } a=L, s'=s \\ 0 & \text{if } a=L, s' \neq s \\ 0.5 & \text{if } a \neq L \end{cases}$$

$$Z(o|a,s') = \begin{cases} 0.5 & \text{if } a \neq L \\ 0.85 & \text{if } a=L \text{ and } o=s' \\ 0.15 & \text{if } a=L \text{ and } o \neq s' \end{cases}$$

$$\gamma = 0.95$$

Hidden Markov Models and Beliefs


Let

Hidden Markov Models and Beliefs




Let

Hidden Markov Models and Beliefs



Let


Hidden Markov Models and Beliefs



$$P(s_{t_{\#}} \mid s_0, a_0, \dots, s_t, a_t) = T(s_{t+1} \mid s_t, a_t)$$

Let


Hidden Markov Models and Beliefs



$$P(s_{t_1} \mid s_0, a_0, \dots, s_t, a_t) = T(s_{t+1} \mid s_t, a_t)$$

Let

Hidden Markov Models and Beliefs




Let

$$P(s_{t_1} \mid s_0, a_0, \dots, s_t, a_t) = T(s_{t+1} \mid s_t, a_t)$$

$$P(s_{t_1} \mid o_0, a_0, \dots, o_t, a_t) = P(s_{t+1} \mid a_t, o_{t+t})????$$

Hidden Markov Models and Beliefs




Let

$$P(s_{t_1} \mid s_0, a_0, \dots, s_t, a_t) = T(s_{t+1} \mid s_t, a_t)$$

~~$$P(s_{t_1} \mid o_0, a_0, \dots, o_t, a_t) = P(s_{t+1} \mid a_t, o_{t+t})????$$~~

Hidden Markov Models and Beliefs




Let

- $b_0(s) \equiv P(s_0 = s)$

$$P(s_{t_1} \mid s_0, a_0, \dots, s_t, a_t) = T(s_{t+1} \mid s_t, a_t)$$

~~$$P(s_{t_1} \mid o_0, a_0, \dots, o_t, a_t) = P(s_{t+1} \mid a_t, o_{t+t})????$$~~

Hidden Markov Models and Beliefs




Let

- $b_0(s) \equiv P(s_0 = s)$
- $h_t \equiv (b_0, a_0, o_1, a_1, \dots, a_{t-1}, o_t)$

$$P(s_{t_1} \mid s_0, a_0, \dots, s_t, a_t) = T(s_{t+1} \mid s_t, a_t)$$

~~$$P(s_{t_1} \mid o_0, a_0, \dots, o_t, a_t) = P(s_{t+1} \mid a_t, o_{t+t})????$$~~

Hidden Markov Models and Beliefs



Let

- $b_0(s) \equiv P(s_0 = s)$
- $h_t \equiv (b_0, a_0, o_1, a_1, \dots, a_{t-1}, o_t)$
- $b_t(s) \equiv P(s_t = s \mid h_t)$

$$P(s_{t_1} \mid s_0, a_0, \dots, s_t, a_t) = T(s_{t+1} \mid s_t, a_t)$$

~~$$P(s_{t_1} \mid o_0, a_0, \dots, o_t, a_t) = P(s_{t+1} \mid a_t, o_{t+t})????$$~~

Bayesian Belief Updates

$$b_t = \tau(b_{t-1}, a_{t-1}, o_t)$$

$$h_t = (h_{t-1}, a_{t-1}, o_t)$$

$$\begin{aligned} b_t &= P(s_t | h_t) = P(s_t | h_{t-1}, a_{t-1}, o_t) \\ &= \frac{P(o_t | s_t, h_{t-1}, a_{t-1}) P(s_t | h_{t-1}, a_{t-1})}{P(o_t | h_{t-1}, a_{t-1})} \end{aligned}$$

$$o_t \perp h_t | s_t, a_{t-1}$$


$$\propto P(o_t | s_t, h_{t-1}, a_{t-1}) P(s_t | h_{t-1}, a_{t-1})$$

$$= P(o_t | a_{t-1}, s_t) \sum_{s_{t-1}} P(s_t | s_{t-1}, a_{t-1}) P(s_{t-1} | h_{t-1})$$

$$b_t(s_t) \propto Z(o_t | a_{t-1}, s_t) \sum_{s_{t-1}} T(s_t | s_{t-1}, a_{t-1}) b_{t-1}(s_{t-1})$$

$$\boxed{b' \propto Z(o | a, s') \sum_s T(s' | s, a) b(s)}$$

$$b' = \tau(b, a, o)$$



Filtering Loop

$$b = b_0$$

loop

receiving o

for $s' \in S$

$$b'(s') \leftarrow Z(o | a, s') \sum_s T(s' | s, a) b(s)$$

$$b' \leftarrow b' / \sum_{s'} b'(s')$$

$$b \leftarrow b'$$

Tiger Example

