

Recap

- Alpha Vectors
- Best solver for discrete POMDPs:

SARSO

POMDP Computational Complexity

POMDP Computational Complexity

Sad facts ● 

- Infinite horizon POMDPs are *undecidable*

POMDP Computational Complexity

Sad facts ●

- Infinite horizon POMDPs are *undecidable*
- Finite horizon POMDPs are *PSPACE Complete*

POMDP Computational Complexity


Sad facts ●

- Infinite horizon POMDPs are *undecidable*
- Finite horizon POMDPs are *PSPACE Complete*
 - Among the hardest problems that can be solved using a polynomial amount of space

POMDP Computational Complexity

Sad facts ●


- Infinite horizon POMDPs are *undecidable*
- Finite horizon POMDPs are *PSPACE Complete*
 - Among the hardest problems that can be solved using a polynomial amount of space



POMDP Computational Complexity

Sad facts ●


- Infinite horizon POMDPs are *undecidable*
- Finite horizon POMDPs are *PSPACE Complete*
 - Among the hardest problems that can be solved using a polynomial amount of space



POMDP Computational Complexity

Sad facts ●


- Infinite horizon POMDPs are *undecidable*
- Finite horizon POMDPs are *PSPACE Complete*
 - Among the hardest problems that can be solved using a polynomial amount of space



POMDP Computational Complexity

Sad facts ●


- Infinite horizon POMDPs are *undecidable*
- Finite horizon POMDPs are *PSPACE Complete*
 - Among the hardest problems that can be solved using a polynomial amount of space
 - Any algorithm that can solve a general POMDP will have exponential complexity



POMDP Computational Complexity

Sad facts ●

- Infinite horizon POMDPs are *undecidable*
- Finite horizon POMDPs are *PSPACE Complete*
 - Among the hardest problems that can be solved using a polynomial amount of space
 - Any algorithm that can solve a general POMDP will have exponential complexity (we think)



Approximate POMDP Solutions

Approximate POMDP Solutions

Numerical Approximations

(approximately solve original problem)

Approximate POMDP Solutions

Numerical Approximations

(approximately solve original problem)



Offline

Approximate POMDP Solutions

Numerical Approximations

(approximately solve original problem)



Offline

Last week

Approximate POMDP Solutions

Numerical Approximations

(approximately solve original problem)



Offline

Last week



Online

Approximate POMDP Solutions

Numerical Approximations

(approximately solve original problem)



Offline

Last week



Online

Thursday

Approximate POMDP Solutions

Numerical Approximations

(approximately solve original problem)



Offline

Last week



Online

Thursday

Formulation Approximations

(solve a slightly different problem)

Approximate POMDP Solutions

Numerical Approximations

(approximately solve original problem)



Offline

Last week



Online


Thursday

Formulation Approximations

(solve a slightly different problem)

Today!

Rotor Failure Example



$$s = (\underbrace{x, y, z, \dot{x}, \dot{y}, \dot{z}, \phi, \Theta, \psi, p, q, r}_{\text{known}}, \underbrace{m_1, m_2, \dots, m_6}_{\text{unknown}})$$

$m \in \{\text{nominal}, \text{failed}\}$

Certainty - Equivalence

POMDP Objective

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

POMDP Objective

$$\pi^* = \underset{\pi: B \rightarrow A}{\operatorname{argmax}} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

Certainty Equivalent

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

Certainty Equivalent

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

$$b' = \tau(b, a, o)$$

Certainty Equivalent

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

$$\pi_{\text{CE}}(b) = \underset{\pi_{\text{MDP}}^*}{\overset{\uparrow}{\pi_s}}(\underset{s \sim b}{\mathbb{E}[s]})$$

$\pi_s(\hat{s})$ $\hat{s} = \mathbb{E}_{s \sim b}[s]$

$$b' = \tau(b, a, o)$$

Certainty Equivalent

$$\begin{aligned} S &= \mathbb{R}^n \\ A &= \mathbb{R}^m \\ O &= \mathbb{R}^p \end{aligned}$$

$$T(s' | s, a) = \mathcal{N}(s' | \mathbf{T}_s s + \mathbf{T}_a a, \Sigma_s)$$

$$O(o | s') = \mathcal{N}(o | \mathbf{O}_s s', \Sigma_o)$$

$$R(s, a) = -s^T R_s s + a^T R_a a$$

Optimal for LQG ← Gaussian
 ↑ ↑
 Linear Quadratic

Kalman Filter

$$b(s) = \mathcal{N}(s | \mu_b, \Sigma_b)$$

Prediction

$$\mu_p \leftarrow \mathbf{T}_s \mu_b + \mathbf{T}_a a$$

$$\Sigma_p \leftarrow \mathbf{T}_s \Sigma_b \mathbf{T}_s^T + \Sigma_s$$

Update

$$\mathbf{K} \leftarrow \Sigma_p \mathbf{O}_s^T (\mathbf{O}_s \Sigma_p \mathbf{O}_s^T + \Sigma_o)^{-1}$$

$$\mu_b \leftarrow \mu_p + \mathbf{K}(\underline{o} - \mathbf{O}_s \mu_p)$$

$$\Sigma_b \leftarrow (\mathbf{I} - \mathbf{K} \mathbf{O}_s) \Sigma_p$$

$$x^*(b) = K_a \text{mean}(b)$$

QMDP

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

QMDP

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

$$b' = \tau(b, a, o)$$

QMDP

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

$(s, A, T, R, O, Z, \gamma)$
solve


optimal Q-value
for the fully
observable
MDP

$$\pi_{\text{QMDP}}(b) = \operatorname{argmax}_{a \in A} \mathbb{E}_{s \sim b} [Q_{\text{MDP}}(s, a)]$$

$\operatorname{argmax}_{a \in A} \pi_s(\mathbb{E}_{s \sim b}[\mathbb{E}[\mathbb{E}]])$ CE

$$b' = \tau(b, a, o)$$

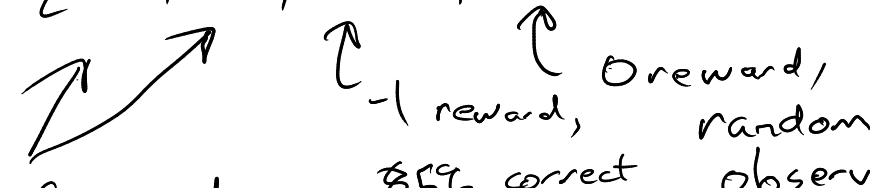
Value



$$\operatorname{argmax}_{a \in A} \alpha_a \cdot b$$

$$\alpha_a[s] = Q_{\text{MDP}}(s, a)$$

Example: Tiger POMDP with Waiting

$A = \{OL, OR, \text{Listen}, \text{Wait}\}$

 $+10$ if good
 -100 if bad
 -1 reward,
 55% correct
 observation
 0 reward,
 random
 observation

$$b(OL) = 50\%$$


optimal = L

QMDP solution?

$$Q_{MDP}(\text{any}, \text{open}) = 10$$


$$\gamma = 0.9$$

pseudo alpha vectors




$$-1 + \gamma \max_a Q_{MDP}(s, a)$$

POMDP Example: Light-Dark



$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ \underbrace{s' = s + a} & & \underbrace{o \sim \mathcal{N}(s, s - 10)} \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\} \\ R(s, a) &= \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}\end{aligned}$$

POMDP Example: Light-Dark




$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\} \\ R(s, a) &= \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}\end{aligned}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations




$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\}\end{aligned}$$

$$R(s, a) = \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations




$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\}\end{aligned}$$

$$R(s, a) = \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations




$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\} \\ R(s, a) &= \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}\end{aligned}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations

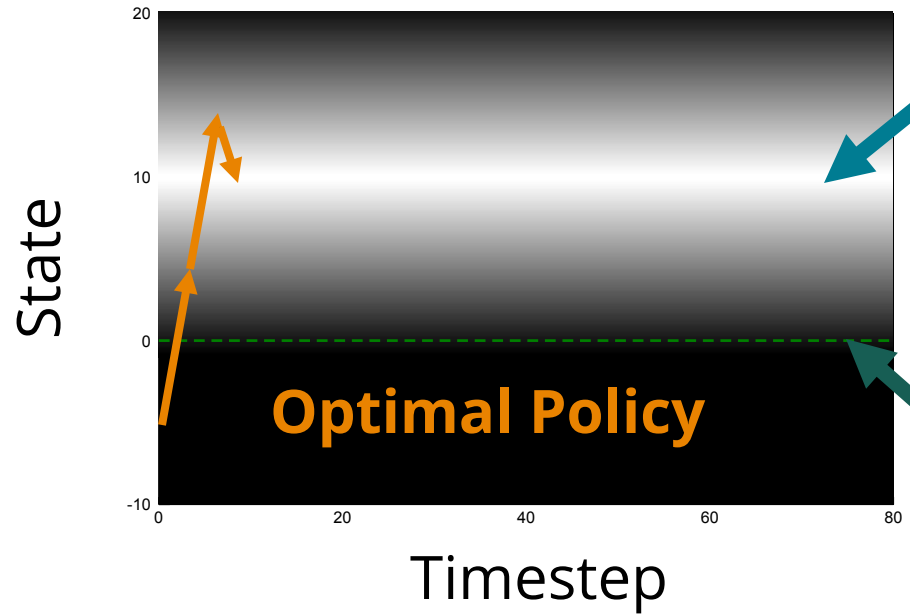


$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\} \\ R(s, a) &= \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}\end{aligned}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations




$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\} \\ R(s, a) &= \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}\end{aligned}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations



$$\mathcal{S} = \mathbb{Z}$$

$$\mathcal{O} = \mathbb{R}$$

$$s' = s + a \quad o \sim \mathcal{N}(s, s - 10)$$


$$\mathcal{A} = \{-10, -1, 0, 1, 10\}$$

$$R(s, a) = \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations




$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\} \\ R(s, a) &= \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}\end{aligned}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations



$$\mathcal{S} = \mathbb{Z}$$

$$\mathcal{O} = \mathbb{R}$$

$$s' = s + a \quad o \sim \mathcal{N}(s, s - 10)$$


$$\mathcal{A} = \{-10, -1, 0, 1, 10\}$$

$$R(s, a) = \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}$$

Goal: $a = 0$ at $s = 0$

POMDP Example: Light-Dark

Accurate Observations




$$\begin{aligned}\mathcal{S} &= \mathbb{Z} & \mathcal{O} &= \mathbb{R} \\ s' &= s + a & o &\sim \mathcal{N}(s, s - 10) \\ \mathcal{A} &= \{-10, -1, 0, 1, 10\}\end{aligned}$$


$$R(s, a) = \begin{cases} 100 & \text{if } a = 0, s = 0 \\ -100 & \text{if } a = 0, s \neq 0 \\ -1 & \text{otherwise} \end{cases}$$

Goal: $a = 0$ at $s = 0$

POMDP Solution



QMDP




Same as **full observability**
on the next step

Information Gathering

QMDP


Full POMDP



Information Gathering


QMDP

Full POMDP

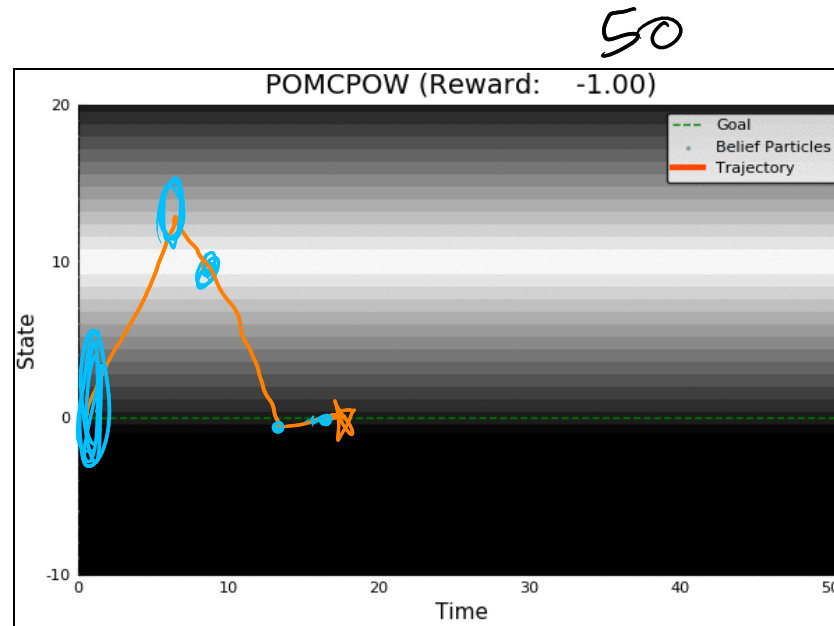


Information Gathering

QMDP




Full POMDP



QMDP


INDUSTRIAL GRADE



QMDP

ACAS X

[Kochenderfer, 2011]



Hindsight Optimization

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$

m scenarios

$$\pi_{\text{Hop}}(b) = \operatorname{argmax}_{a_{0:\infty}^i} \frac{1}{m} \sum_{i=1}^m \gamma^T R(s_T^i, a_T^i)$$

Subject to

$$s_{t+1}^i = G(s_t^i, a_t^i, \phi_t^i)$$

fixed outcome

$$a_t^i = a_t^j \quad \forall i, j$$

\mathcal{A}

FIB

POMDP Objective

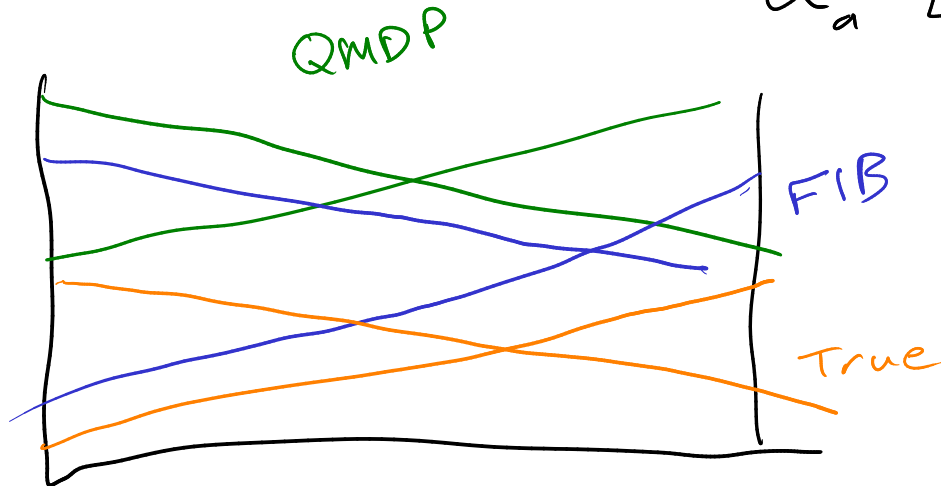
$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$\pi_{\text{FIB}}(b) = \operatorname{argmax}_{a \in A} \alpha_a^T b$$

$$b' = \tau(b, a, o)$$

iterate

$$\alpha_a^{(k+1)}[s] = R(s, a) + \gamma \sum_o \max_{a'} \sum_{s'} Z(o|a, s') T(s'|s, a) \alpha_{a'}^{(k)}[s']$$



k-Markov


POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(b_t)) \right]$$

$$b' = \tau(b, a, o)$$



← previous observations



Open Loop

POMDP Objective

$$\pi^* = \operatorname{argmax}_{\pi: B \rightarrow A} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(\underline{b_t})) \right]$$

$$b' = \tau(b, a, o)$$

$$\max_{a_1, a_2, \dots, a_T} \mathbb{E} \left[\sum_{t=0}^T \gamma^t R(s_t, a_t) \right]$$