


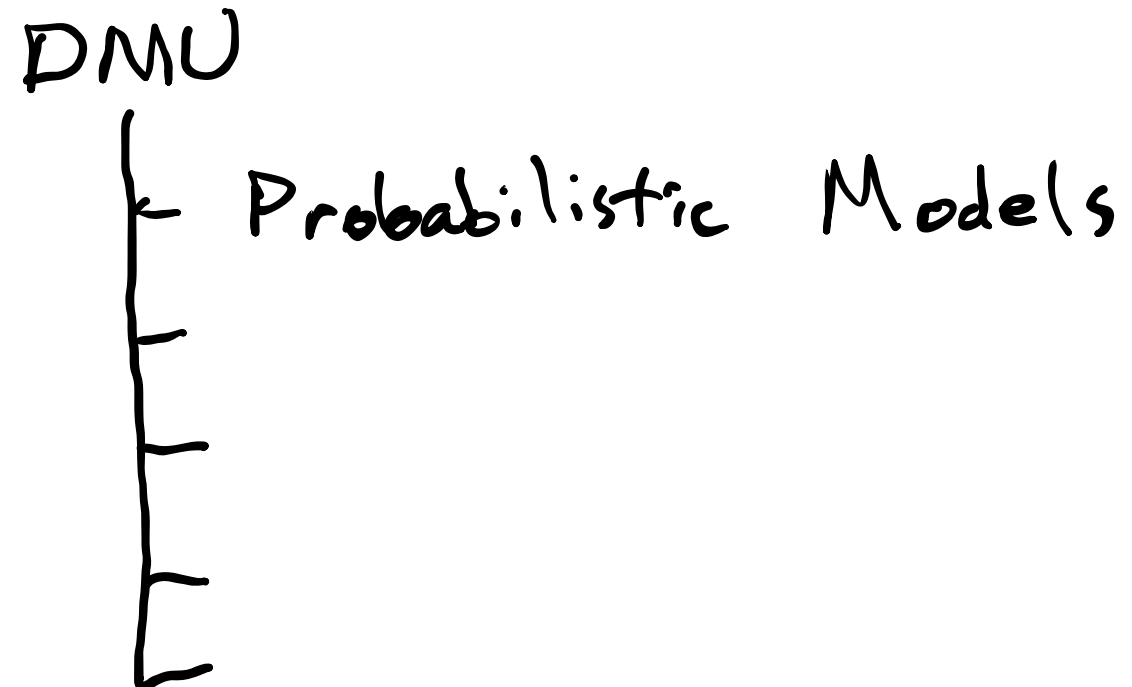
# Recap

# Recap



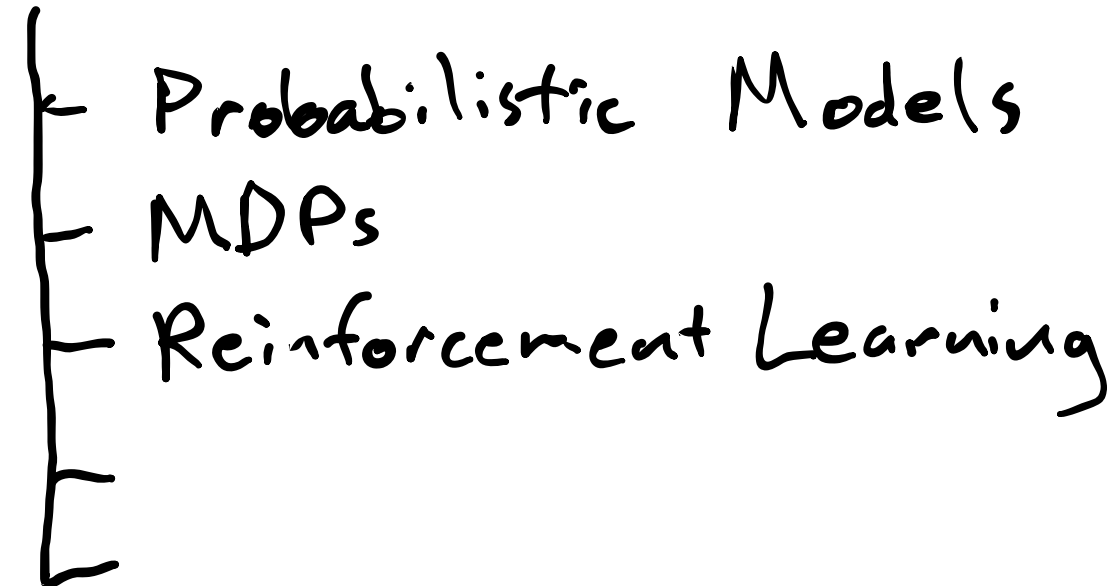
# Recap

DMU



# Recap

DMU



# Recap

DMU

- Probabilistic Models
- MDPs
- Reinforcement Learning
- POMDPs
- Games

# Probabilistic Models

# Probabilistic Models

3 Rules

$P(A)$   
 $P(A, B)$   
 $P(A|B)$

# Probabilistic Models

## 3 Rules

$$\begin{aligned} P(A) \\ P(A, B) \\ P(A|B) \end{aligned}$$

$$\begin{aligned} 1. \quad & 0 \leq P(X | Y) \leq 1 \\ & \sum_{x \in X} P(x | Y) = 1 \end{aligned}$$



# Probabilistic Models

$P(A)$   
 $P(A, B)$   
 $P(A|B)$

## 3 Rules

1.  $0 \leq P(X | Y) \leq 1$

$$\sum_{x \in X} P(x | Y) = 1$$

2.  $P(X) = \sum_{y \in Y} P(X, y)$

# Probabilistic Models

$$P(A)$$
$$P(A, B)$$
$$P(A|B)$$

## 3 Rules

$$1. 0 \leq P(X | Y) \leq 1$$

$$\sum_{x \in X} P(x | Y) = 1$$

$$2. P(X) = \sum_{y \in Y} P(X, y)$$

$$3. P(X | Y) = \frac{P(X, Y)}{P(Y)}$$

# Probabilistic Models

$$\begin{aligned} P(A) \\ P(A, B) \\ P(A|B) \end{aligned}$$

## 3 Rules

$$1. 0 \leq P(X | Y) \leq 1$$

$$\sum_{x \in X} P(x | Y) = 1$$

$$2. P(X) = \sum_{y \in Y} P(X, y)$$

$$3. P(X | Y) = \frac{P(X, Y)}{P(Y)}$$

## Bayes Rule

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

# Probabilistic Models

$$\begin{aligned} P(A) \\ P(A, B) \\ P(A|B) \end{aligned}$$

## 3 Rules

$$1. 0 \leq P(X | Y) \leq 1$$

$$\sum_{x \in X} P(x | Y) = 1$$

$$2. P(X) = \sum_{y \in Y} P(X, y)$$

$$3. P(X | Y) = \frac{P(X, Y)}{P(Y)}$$

## Bayes Rule

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$


## Independence

$$A \perp B \iff P(A, B) = P(A)P(B)$$


$$A \perp B | C \iff P(A, B | C) = P(A | C)P(B | C)$$

# Bayesian Networks

# Bayesian Networks




# Bayesian Networks



$$P(X_i \mid X_{1:n \setminus i}) \stackrel{?}{=} P(X_i \mid Pa(X_i))$$

# Bayesian Networks




$$\cancel{P(X_i \mid X_{1:n \setminus i}) \stackrel{?}{=} P(X_i \mid Pa(X_i))}$$

No!



# Bayesian Networks




$$\cancel{P(X_i \mid X_{1:n \setminus i}) \stackrel{?}{=} P(X_i \mid Pa(X_i))}$$

No!

**Chain Rule**

$$P(X_{1:n}) = \prod_i P(X_i \mid Pa(X_i))$$

# Bayesian Networks



$$\cancel{P(X_i \mid X_{1:n \setminus i}) \stackrel{?}{=} P(X_i \mid Pa(X_i))}$$

No!


## Chain Rule

$$P(X_{1:n}) = \prod_i P(X_i \mid Pa(X_i))$$

## Sampling

Topological sort, then  
sample from each node

# Bayesian Networks



$$\cancel{P(X_i | X_{1:n \setminus i}) \stackrel{?}{=} P(X_i | Pa(X_i))}$$

No!

## Sampling

Topological sort, then  
sample from each node


## Chain Rule

$$P(X_{1:n}) = \prod_i P(X_i | Pa(X_i))$$

## Conditional Independence

$X \perp Y \mid \mathcal{C}$  if all paths between  $X$  and  $Y$  are d-separated by  $\mathcal{C}$

# Bayesian Networks



$$\cancel{P(X_i | X_{1:n \setminus i}) \stackrel{?}{=} P(X_i | Pa(X_i))}$$

No!

## Chain Rule

$$P(X_{1:n}) = \prod_i P(X_i | Pa(X_i))$$

## Sampling

Topological sort, then  
sample from each node


## Conditional Independence

$X \perp Y \mid \mathcal{C}$  if all paths between  $X$  and  $Y$  are d-separated by  $\mathcal{C}$

## Inference

- Input: BN, evidence values
- Output: Distribution of targets

# Bayesian Networks



$$\cancel{P(X_i | X_{1:n \setminus i}) \stackrel{?}{=} P(X_i | Pa(X_i))}$$

No!

## Chain Rule

$$P(X_{1:n}) = \prod_i P(X_i | Pa(X_i))$$

## Sampling

Topological sort, then  
sample from each node

## Conditional Independence

$X \perp Y \mid \mathcal{C}$  if all paths between  $X$  and  $Y$  are d-separated by  $\mathcal{C}$


## Inference

- Input: BN, evidence values
- Output: Distribution of targets

Exact: NP-Hard

Approximate via sampling: Direct,  
Likelihood Weighted, Gibbs

# Bayesian Networks



$$\cancel{P(X_i | X_{1:n \setminus i}) \stackrel{?}{=} P(X_i | Pa(X_i))}$$

No!

## Chain Rule

$$P(X_{1:n}) = \prod_i P(X_i | Pa(X_i))$$

## Sampling

Topological sort, then  
sample from each node

## Conditional Independence

$X \perp Y \mid \mathcal{C}$  if all paths between  $X$  and  $Y$  are d-separated by  $\mathcal{C}$

## Inference

- Input: BN, evidence values
- Output: Distribution of targets

Exact: NP-Hard

Approximate via sampling: Direct,  
Likelihood Weighted, Gibbs

## Learning

- Input: Data
- Output: BN structure and parameters

# Markov Decision Processes

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$



# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

$$\underset{\pi}{\text{maximize}} \quad E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

$$\underset{\pi}{\text{maximize}} \ E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

$$Q^{\pi}(s, a) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a, a_t = \pi(s_t) \right]$$

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

$$\underset{\pi}{\text{maximize}} \ E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

$$Q^{\pi}(s, a) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a, a_t = \pi(s_t) \right]$$

$$V^{\pi}(s) = \max_a Q^{\pi}(s, a)$$

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

$$\underset{\pi}{\text{maximize}} \ E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

$$Q^{\pi}(s, a) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a, a_t = \pi(s_t) \right]$$

$$V^{\pi}(s) = \max_a Q^{\pi}(s, a)$$

$$V^{\pi}(s) = R(s, a) + \gamma E[V^{\pi}(s')]$$

$$V^*(s) = \max_a \{R(s, a) + \gamma E[V^*(s')]\}$$

$$B[V](s) = \max_a \{R(s, a) + \gamma E[V(s')]\}$$

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

$$\underset{\pi}{\text{maximize}} \ E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

$$Q^{\pi}(s, a) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a, a_t = \pi(s_t) \right]$$

$$V^{\pi}(s) = \max_a Q^{\pi}(s, a)$$

$$V^{\pi}(s) = R(s, a) + \gamma E[V^{\pi}(s')]$$

Policy Evaluation

$$V^*(s) = \max_a \{R(s, a) + \gamma E[V^*(s')]\}$$

$$B[V](s) = \max_a \{R(s, a) + \gamma E[V(s')]\}$$

# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

$$\underset{\pi}{\text{maximize}} \ E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

$$Q^{\pi}(s, a) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a, a_t = \pi(s_t) \right]$$

$$V^{\pi}(s) = \max_a Q^{\pi}(s, a)$$

$$V^{\pi}(s) = R(s, a) + \gamma E[V^{\pi}(s')]$$

Policy Evaluation

$$V^*(s) = \max_a \{R(s, a) + \gamma E[V^*(s')]\}$$

Bellman's Equation: Certificate of Optimality

$$B[V](s) = \max_a \{R(s, a) + \gamma E[V(s')]\}$$



# Markov Decision Processes

$$(S, A, R, T, \gamma)$$

Examples:  $S = \{1, 2, 3\}$  or  $S = \mathbb{R}^2$

$$s = (x, \dot{x}) \in S = \mathbb{R}^2$$

$$\underset{\pi}{\text{maximize}} \ E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

$$Q^{\pi}(s, a) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a, a_t = \pi(s_t) \right]$$

$$V^{\pi}(s) = \max_a Q^{\pi}(s, a)$$

$$V^{\pi}(s) = R(s, a) + \gamma E[V^{\pi}(s')]$$

Policy Evaluation

$$V^*(s) = \max_a \{R(s, a) + \gamma E[V^*(s')]\}$$

Bellman's Equation: Certificate of Optimality

$$B[V](s) = \max_a \{R(s, a) + \gamma E[V(s')]\}$$

Bellman's Operator

# Offline MDP Algorithms

# Offline MDP Algorithms

## **Policy Iteration**

loop

Evaluate Policy

Improve Policy

# Offline MDP Algorithms

## Policy Iteration

loop

Evaluate Policy

Improve Policy

## Value Iteration

loop

$$V \leftarrow B[V]$$

# Offline MDP Algorithms

## Policy Iteration

loop

Evaluate Policy

Improve Policy

Converges because  
policy always improves  
and there are a finite  
number of policies

## Value Iteration

loop

$$V \leftarrow B[V]$$

# Offline MDP Algorithms

## Policy Iteration

loop

Evaluate Policy

Improve Policy

Converges because  
policy always improves  
and there are a finite  
number of policies

## Value Iteration

loop

$$V \leftarrow B[V]$$

Converges because  $B$  is  
a contraction mapping

# Online MDP Planning

# Online MDP Planning

**Monte Carlo Tree Search**



# Online MDP Planning

## Monte Carlo Tree Search

Search

Expand

Rollout

Backup

# Online MDP Planning

## Monte Carlo Tree Search

Search  
Expand  
Rollout  
Backup


$$Q(s,a) + c \sqrt{\frac{\log N(s)}{N(s,a)}}$$

# Online MDP Planning

## Monte Carlo Tree Search

Search  
Expand  
Rollout  
Backup

$$Q(s,a) + c \sqrt{\frac{\log N(s)}{N(s,a)}}$$




# Online MDP Planning

## Monte Carlo Tree Search

Search  
Expand  
Rollout  
Backup

$$Q(s,a) + c \sqrt{\frac{\log N(s)}{N(s,a)}}$$




## Sparse Sampling

# Online MDP Planning



## Monte Carlo Tree Search

Search  
Expand  
Rollout  
Backup

$$Q(s,a) + c \sqrt{\frac{\log N(s)}{N(s,a)}}$$



## Sparse Sampling




# Online MDP Planning


## Monte Carlo Tree Search

Search  
Expand  
Rollout  
Backup

$$Q(s,a) + c \sqrt{\frac{\log N(s)}{N(s,a)}}$$




## Sparse Sampling



Guarantees *independent* of  $|S|!!$

# Reinforcement Learning



Challenges:

1. Exploration and Exploitation
2. Credit Assignment
3. Generalization

# Exploration



# Exploration


## Bandits

- $\epsilon$ -greedy
- softmax
- UCB
- Thompson Sampling
- Optimal DP Solution (solving a POMDP!)

# Exploration

## Bandits

- $\epsilon$ -greedy
- softmax
- UCB
- Thompson Sampling
- Optimal DP Solution (solving a POMDP!)




Montezuma's Revenge!

# Exploration

## Bandits

- $\epsilon$ -greedy
- softmax
- UCB
- Thompson Sampling
- Optimal DP Solution (solving a POMDP!)

- Pseudocounts




Montezuma's Revenge!

# Exploration

## Bandits

- $\epsilon$ -greedy
- softmax
- UCB
- Thompson Sampling
- Optimal DP Solution (solving a POMDP!)




Montezuma's Revenge!

- Pseudocounts
- Curiosity: extra reward for bad prediction

# Exploration

## Bandits

- $\epsilon$ -greedy
- softmax
- UCB
- Thompson Sampling
- Optimal DP Solution (solving a POMDP!)



Montezuma's Revenge!

- Pseudocounts
- Curiosity: extra reward for bad prediction
- Random network distillation

# RL Algorithms

# RL Algorithms

Model  
Based


Model  
Free



# RL Algorithms

Model  
Based

Model  
Free





# RL Algorithms

Model  
Based

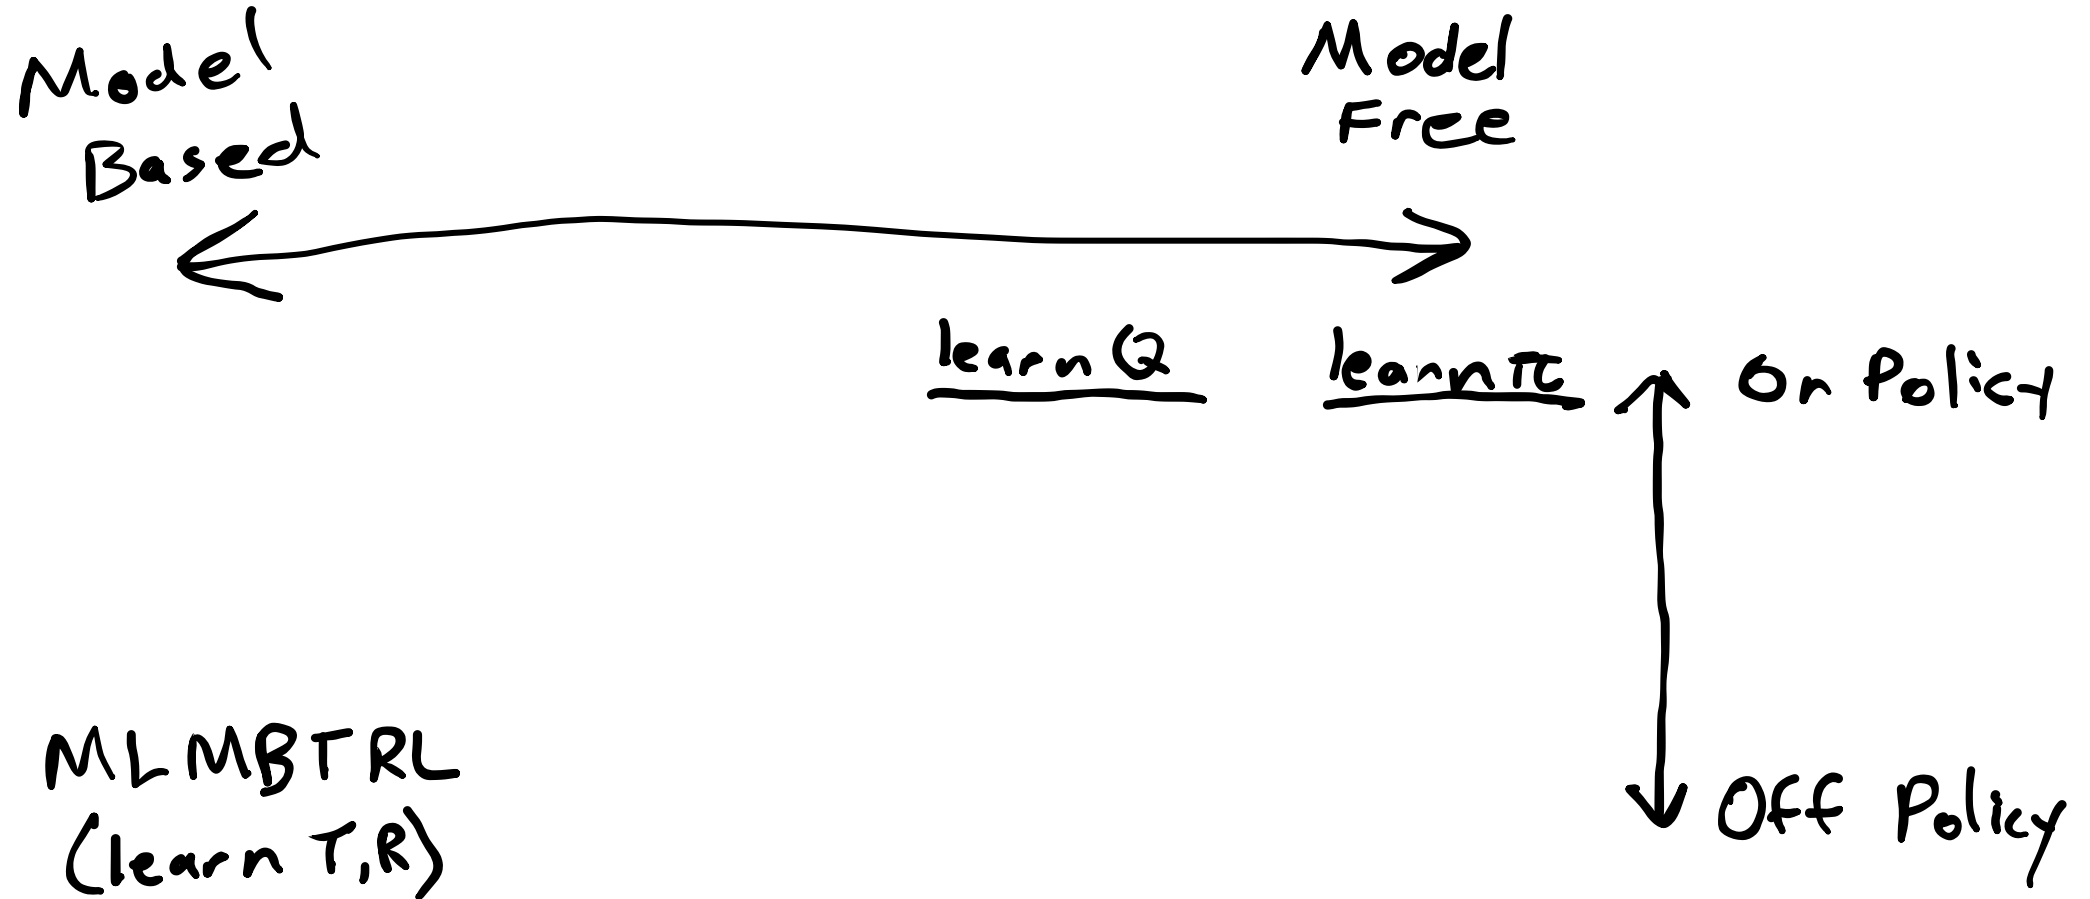
Model  
Free




MLMBTRL  
(learn  $T, R$ )

On Policy  
↓  
Off Policy

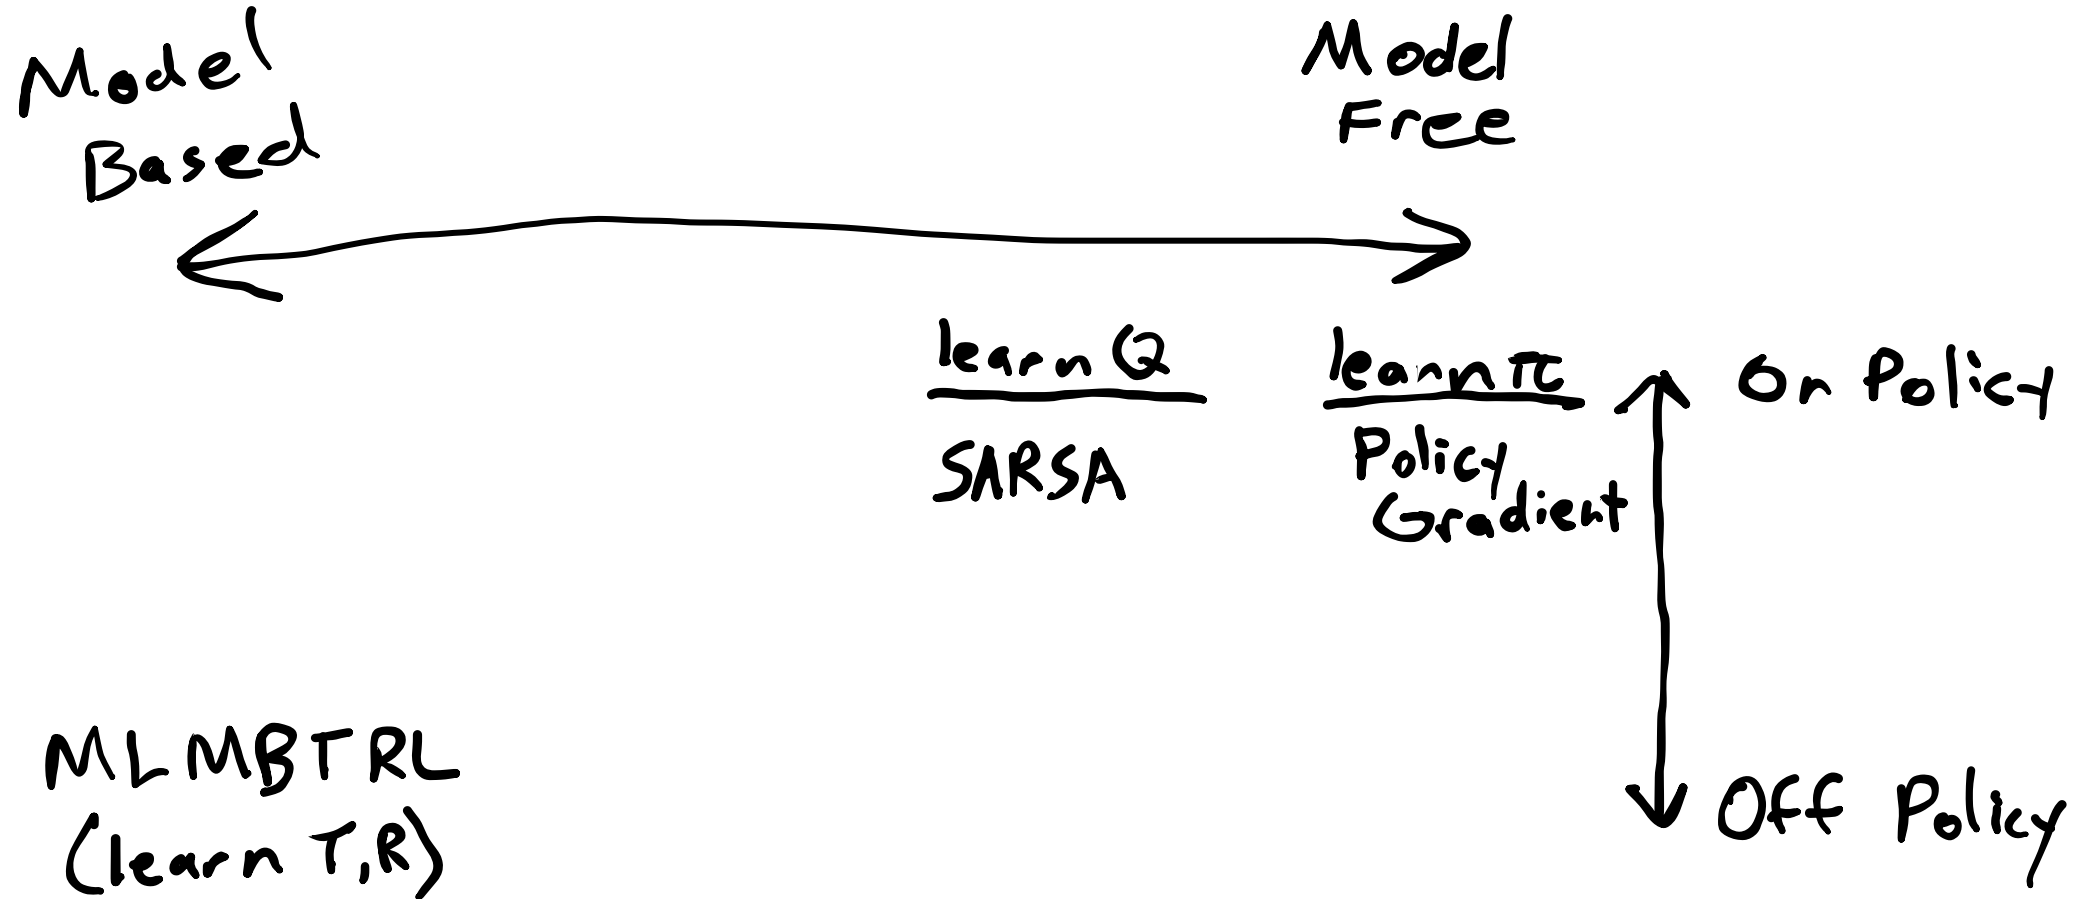
# RL Algorithms




# RL Algorithms



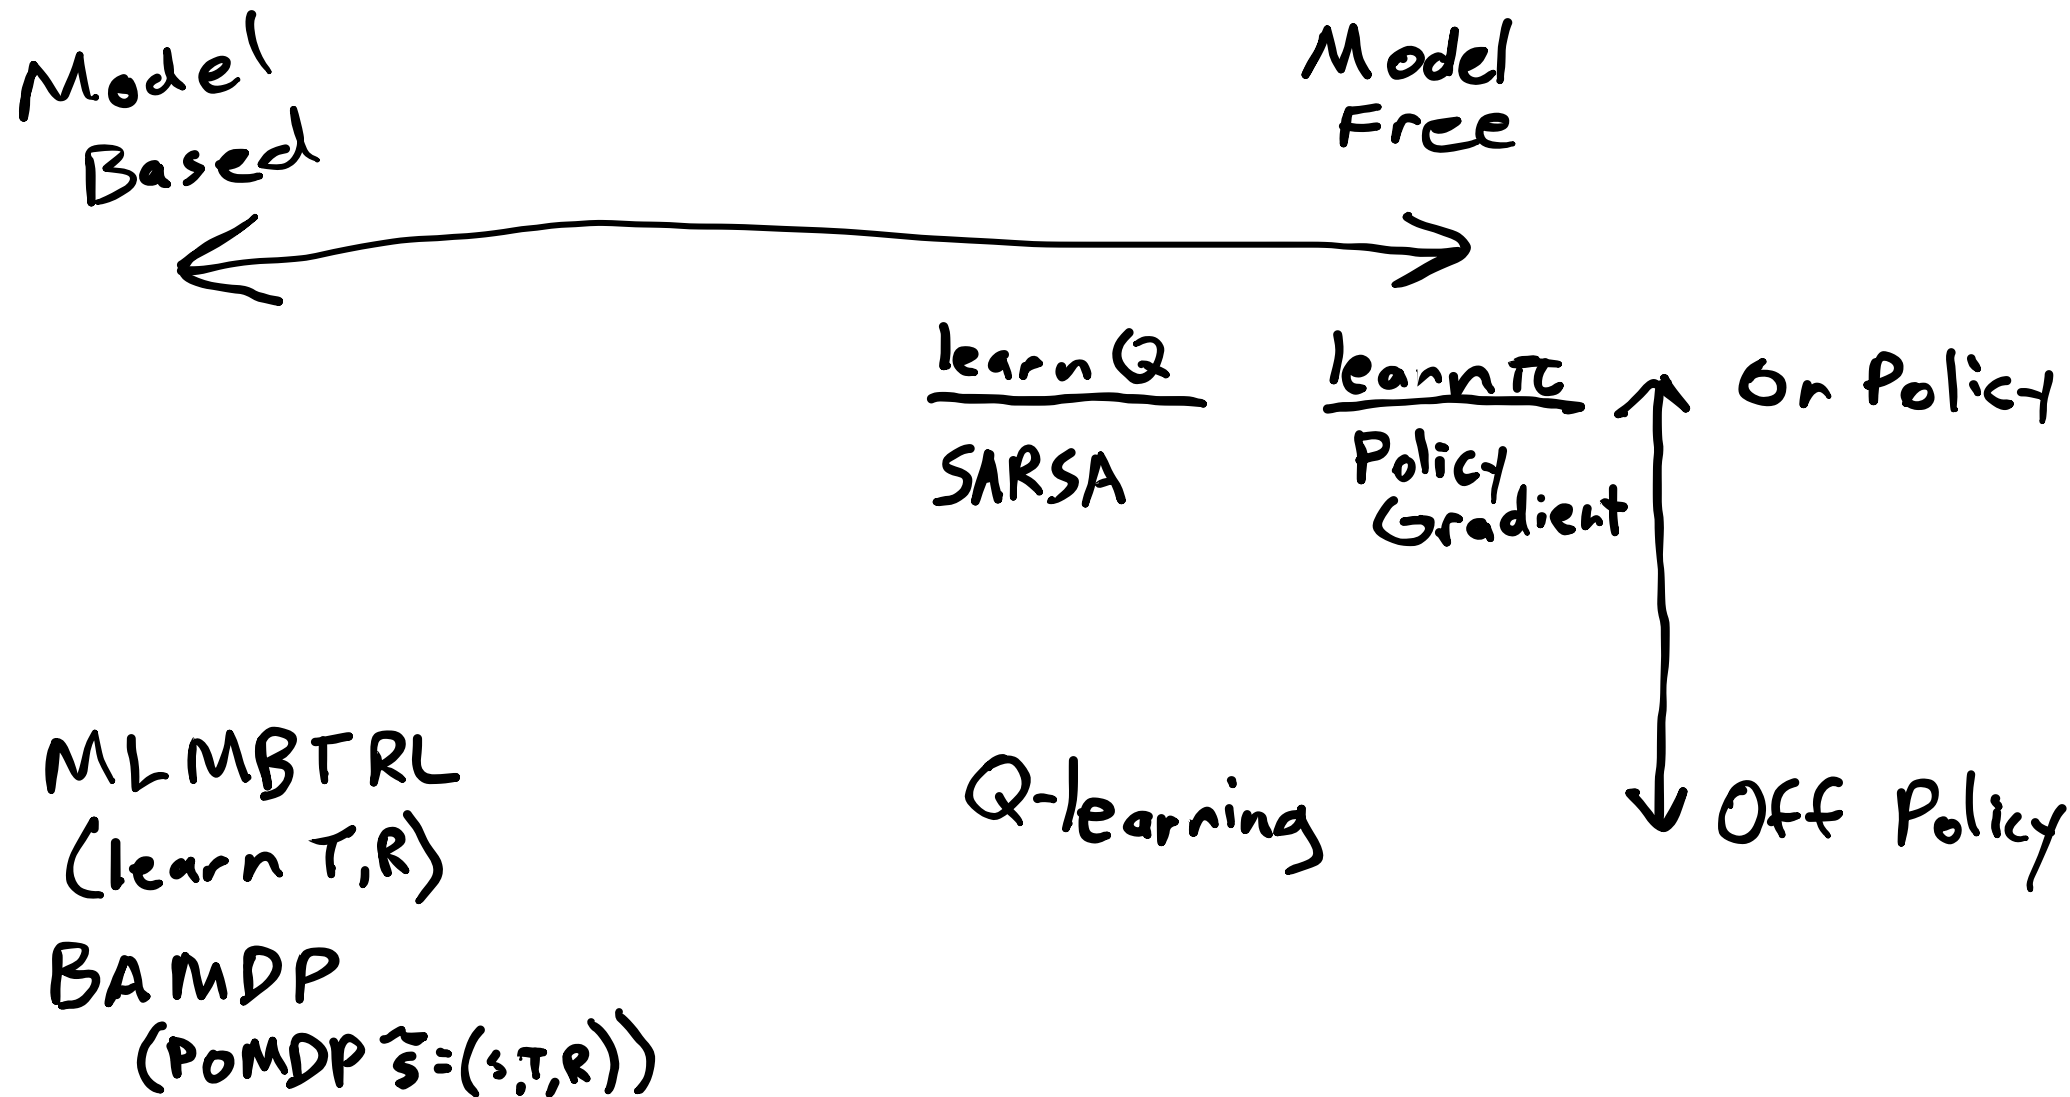
# RL Algorithms



# RL Algorithms



# RL Algorithms



# Policy Gradient

# Policy Gradient

- Likelihood ratio trick



# Policy Gradient

- Likelihood ratio trick

$$\nabla_{\theta} p_{\theta}(\tau) = p_{\theta}(\tau) \log p_{\theta}(\tau)$$

# Policy Gradient

- Likelihood ratio trick
- Causality

$$\nabla_{\theta} p_{\theta}(\tau) = p_{\theta}(\tau) \log p_{\theta}(\tau)$$

# Policy Gradient

- Likelihood ratio trick
- Causality
- Baseline Subtraction

$$\nabla_{\theta} p_{\theta}(\tau) = p_{\theta}(\tau) \log p_{\theta}(\tau)$$

# Policy Gradient

- Likelihood ratio trick
- Causality
- Baseline Subtraction

$$\nabla_{\theta} p_{\theta}(\tau) = p_{\theta}(\tau) \log p_{\theta}(\tau)$$

$$\nabla U(\theta) = \mathbb{E}_{\tau} \left[ \sum_{k=1}^d \nabla_{\theta} \log \pi_{\theta}(a^{(k)} | s^{(k)}) \gamma^{k-1} \left( r_{\text{to-go}}^{(k)} - r_{\text{base}}(s^{(k)}) \right) \right]$$

# Policy Gradient

- Likelihood ratio trick
- Causality
- Baseline Subtraction

$$\nabla_{\theta} p_{\theta}(\tau) = p_{\theta}(\tau) \log p_{\theta}(\tau)$$

$$\nabla U(\theta) = \mathbb{E}_{\tau} \left[ \sum_{k=1}^d \nabla_{\theta} \log \pi_{\theta}(a^{(k)} | s^{(k)}) \gamma^{k-1} \left( r_{\text{to-go}}^{(k)} - r_{\text{base}}(s^{(k)}) \right) \right]$$

- Natural Gradient


# Policy Gradient

- Likelihood ratio trick
- Causality
- Baseline Subtraction

$$\nabla_{\theta} p_{\theta}(\tau) = p_{\theta}(\tau) \log p_{\theta}(\tau)$$

$$\nabla U(\theta) = \mathbb{E}_{\tau} \left[ \sum_{k=1}^d \nabla_{\theta} \log \pi_{\theta}(a^{(k)} | s^{(k)}) \gamma^{k-1} (r_{\text{to-go}}^{(k)} - r_{\text{base}}(s^{(k)})) \right]$$

- Natural Gradient



KL div.  
Bound

# Q-Learning

# Q-Learning

## **SARSA**

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r_t + \gamma Q(s', a') - Q(s, a))$$



# Q-Learning

## **SARSA**

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r_t + \gamma Q(s', a') - Q(s, a))$$

Eligibility Traces

# Q-Learning

## **SARSA**

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r_t + \gamma Q(s', a') - Q(s, a))$$

Eligibility Traces

## **Q-learning**

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r_t + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

# Q-Learning

## **SARSA**

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r_t + \gamma Q(s', a') - Q(s, a))$$

Eligibility Traces


## **Q-learning**

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r_t + \gamma \max_{a'} Q(s', a') - Q(s, a))$$


Double Q Learning

# Neural Networks and DQN

# Neural Networks and DQN




# Neural Networks and DQN



$$f_{\theta}(x) = \sigma(W_2\sigma(W_1x + b_1) + b_2)$$


# Neural Networks and DQN



$$f_{\theta}(x) = \sigma(W_2\sigma(W_1x + b_1) + b_2)$$


Backprop

# Neural Networks and DQN



$$f_{\theta}(x) = \sigma(W_2\sigma(W_1x + b_1) + b_2)$$

Backprop





# Actor-Critic

- Actor:  $\pi_{\theta}$
- Critic:  $Q_{\phi}$

## Soft Actor Critic

# Actor-Critic

- Actor:  $\pi_\theta$
- Critic:  $Q_\phi$

## Soft Actor Critic

$$J(\pi) = E \left[ \sum_{t=0}^{\infty} \gamma^t (r_t + \alpha \mathcal{H}(\pi(\cdot | s_t))) \right]$$


# POMDPs

# POMDPs

$$(S, A, T, R, O, Z, \gamma)$$

# POMDPs

$(S, A, T, R, O, Z, \gamma)$




# POMDPs

$(S, A, T, R, O, Z, \gamma)$


## Belief Updates

- Discrete Bayesian Filter
- Particle Filter



# POMDPs


$(S, A, T, R, O, Z, \gamma)$



## Belief Updates


- Discrete Bayesian Filter
- Particle Filter

## Alpha Vectors



# POMDPs

$$(S, A, T, R, O, Z, \gamma)$$




- Each alpha vector corresponds to a conditional plan

## Belief Updates

- Discrete Bayesian Filter
- Particle Filter


## Alpha Vectors





# POMDPs

$$(S, A, T, R, O, Z, \gamma)$$




- Each alpha vector corresponds to a conditional plan
- You can prune alpha vectors by solving an LP

## Belief Updates

- Discrete Bayesian Filter
- Particle Filter

## Alpha Vectors



# POMDP Approximations

# POMDP Approximations

## Formulation

- Certainty Equivalence
- QMDP

# POMDP Approximations

## Formulation

- Certainty Equivalence
- QMDP

## Numerical

# POMDP Approximations


## Formulation

- Certainty Equivalence
- QMDP

## Numerical

### Offline

- Point-Based Value Iteration
- SARSOP



# POMDP Approximations


## Formulation

- Certainty Equivalence
- QMDP

## Numerical

### Offline

- Point-Based Value Iteration
- SARSOP



### Online

- POMCP
- DESPOT

# POMDP Approximations


## Formulation

- Certainty Equivalence
- QMDP

## Numerical


### Offline

- Point-Based Value Iteration
- SARSOP



### Online

- POMCP
- DESPOT



# Simple Games

- ~~Optimal Solutions~~ No!
- Equilibria (e.g. Nash Equilibria)


-1,-1	-3,0
0,-3	-2,-2

- Every Game has at least 1 Nash Equilibrium
- Might be Pure or Mixed




# Turn Taking Games

# Turn Taking Games




# Turn Taking Games




- Value Function Backup

# Turn Taking Games




- Value Function Backup
- $\alpha\beta$  Pruning

# Turn Taking Games



- Value Function Backup
- $\alpha\beta$  Pruning
- Incomplete Information Extensive Form

# Turn Taking Games



- Value Function Backup
- $\alpha\beta$  Pruning
- Incomplete Information Extensive Form

# Recap

# Recap

## Big Problems

1. Immediate and Future Rewards
2. Unknown Models
3. Partial Observability
4. Other Agents

$Q(s,a)$   
RL  
POMDP  
Games.